

Oracle Database 12c: SQL Workshop II

Activity Guide

D80194GC11

Edition 1.1

July 2014

D87272

ORACLE®

Copyright © 2014, Oracle and/or its affiliates. All rights reserved.

Disclaimer

This document contains proprietary information and is protected by copyright and other intellectual property laws. You may copy and print this document solely for your own use in an Oracle training course. The document may not be modified or altered in any way. Except where your use constitutes "fair use" under copyright law, you may not use, share, download, upload, copy, print, display, perform, reproduce, publish, license, post, transmit, or distribute this document in whole or in part without the express authorization of Oracle.

The information contained in this document is subject to change without notice. If you find any problems in the document, please report them in writing to: Oracle University, 500 Oracle Parkway, Redwood Shores, California 94065 USA. This document is not warranted to be error-free.

Restricted Rights Notice

If this documentation is delivered to the United States Government or anyone using the documentation on behalf of the United States Government, the following notice is applicable:

U.S. GOVERNMENT RIGHTS

The U.S. Government's rights to use, modify, reproduce, release, perform, display, or disclose these training materials are restricted by the terms of the applicable Oracle license agreement and/or the applicable U.S. Government contract.

Trademark Notice

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Author

Dimpi Rani Sarmah

Technical Contributors and Reviewers

Nancy Greenberg, Swarnapriya Shridhar, Bryan Roberts, Laszlo Czinkoczki, KimSeong Loh, Brent Dayley, Jim Spiller, Christopher Wensley, Maheshwari Krishnamurthy, Daniel Milne, Michael Almeida, Diganta Choudhury, Manish Pawar, Clair Bennett, Yanti Chang, Joel Goodman, Gerlinde Frenzen, Madhavi Siddireddy

Editors

Raj Kumar, Malavika Jinka

Publishers

Jobi Varghese, Pavithran Adka

This book was published using: **Oracle Tutor**

Table of Contents

Practices for Lesson 1: Introduction	1-1
Practices for Lesson 1: Overview	1-2
Practice 1-1: Using SQL Developer	1-3
Solution 1-1: Using SQL Developer	1-4
Practices for Lesson 2: Introduction to Data Dictionary Views	2-1
Practices for Lesson 2: Overview	2-2
Practice 2-1: Introduction to Data Dictionary Views	2-3
Solution 2-1: Introduction to Data Dictionary Views	2-6
Practices for Lesson 3: Creating Sequences, Synonyms, and Indexes	3-1
Practices for Lesson 3: Overview	3-2
Practice 3-1: Creating Sequences, Synonyms, and Indexes	3-3
Solution 3-1: Creating Sequences, Synonyms, and Indexes	3-5
Practices for Lesson 4: Creating Views	4-1
Practices for Lesson 4: Overview	4-2
Practice 4-1: Creating Views	4-3
Solution 4-1: Creating Views	4-6
Practices for Lesson 5: Managing Schema Objects	5-1
Practices for Lesson 5: Overview	5-2
Practice 5-1: Managing Schema Objects	5-3
Solution 5-1: Managing Schema Objects	5-8
Practices for Lesson 6: Retrieving Data by Using Subqueries	6-1
Practices for Lesson 6: Overview	6-2
Practice 6-1: Retrieving Data by Using Subqueries	6-3
Solution 6-1: Retrieving Data by Using Subqueries	6-8
Practices for Lesson 7: Manipulating Data by Using Subqueries	7-1
Practices for Lesson 7: Overview	7-2
Practice 7-1: Manipulating Data by Using Subqueries	7-3
Solution 7-1: Manipulating Data by Using Subqueries	7-5
Practices for Lesson 8: Controlling User Access	8-1
Practices for Lesson 8: Overview	8-2
Practice 8-1: Controlling User Access	8-3
Solution 8-1: Controlling User Access	8-6
Practices for Lesson 9: Manipulating Data	9-1
Practices for Lesson 9: Overview	9-2
Practice 9-1: Manipulating Data	9-3
Solution 9-1: Manipulating Data	9-8
Practices for Lesson 10: Managing Data in Different Time Zones	10-1
Practices for Lesson 10: Overview	10-2
Practice 10-1: Managing Data in Different Time Zones	10-3
Solution 10-1: Managing Data in Different Time Zones	10-6
Additional Practices and Solutions	11-1
Additional Practices and Solutions	11-2
Additional Practices	11-3

Additional Practices Solutions	11-9
Additional Practices: Case Study	11-15
Additional Practices Solution: Case Study	11-18

Youssef MAZOUZ (z.m.malak@hotmail.com) has a non-transferable
license to use this Student Guide.

Practices for Lesson 1: Introduction

Chapter 1

Practices for Lesson 1: Overview

Practice Overview

You also learn about your user account that you use in this course. You then start SQL Developer, create a new database connection, and browse your HR tables. You also set some SQL Developer preferences, execute SQL statements, and execute an anonymous PL/SQL block by using SQL Worksheet.

Practice 1-1: Using SQL Developer


Tasks

1. Start SQL Developer by using the desktop icon.
2. Create a database connection using the following information:
 - Connection Name: `myconnection`
 - Username: `ora21`
 - Password: `ora21`
 - Hostname: `localhost`
 - Port: `1521`
 - SID: `orcl` (or the value provided to you by the instructor)
3. Test the new connection. If the status is Success, connect to the database by using this new connection.
 - a. Click the Test button in the New/Select Database Connection window.
 - b. If the status is Success, click the Connect button.
4. Browse the structure of the `EMPLOYEES` table and display its data.
 - a. Expand the `myconnection` connection by clicking the plus sign next to it.
 - b. Expand the Tables icon by clicking the plus sign next to it.
 - c. Display the structure of the `EMPLOYEES` table.
 - d. View the data of the `DEPARTMENTS` table.
5. Execute some basic `SELECT` statements to query the data in the `EMPLOYEES` table in the SQL Worksheet area. Use both the Execute Statement (or press F9) and the Run Script (or press F5) icons to execute the `SELECT` statements. Review the results of both methods of executing the `SELECT` statements on the appropriate tabbed pages.
 - a. Write a query to select the last name and salary for any employee whose salary is less than or equal to \$3,000.
 - b. Write a query to display last name, job ID, and commission for all employees who are not entitled to receive a commission.
6. Set your script pathing preference to `/home/oracle/labs/sql2`.
 - a. Select Tools > Preferences > Database > Worksheet.
 - b. Enter the value in the "Select default path to look for scripts" field.
7. Enter the following in the Enter SQL Statement box.

```
SELECT employee_id, first_name, last_name
      FROM employees;
```
8. Save the SQL statement to a script file by using the File > Save menu item.
 - a. Select File > Save.
 - b. Name the file `intro_test.sql`.
 - c. Place the file under your `/home/oracle/labs/sql2/labs` folder.
9. Open and run `confidence.sql` from your `/home/oracle/labs/sql2/labs` folder, and observe the output.

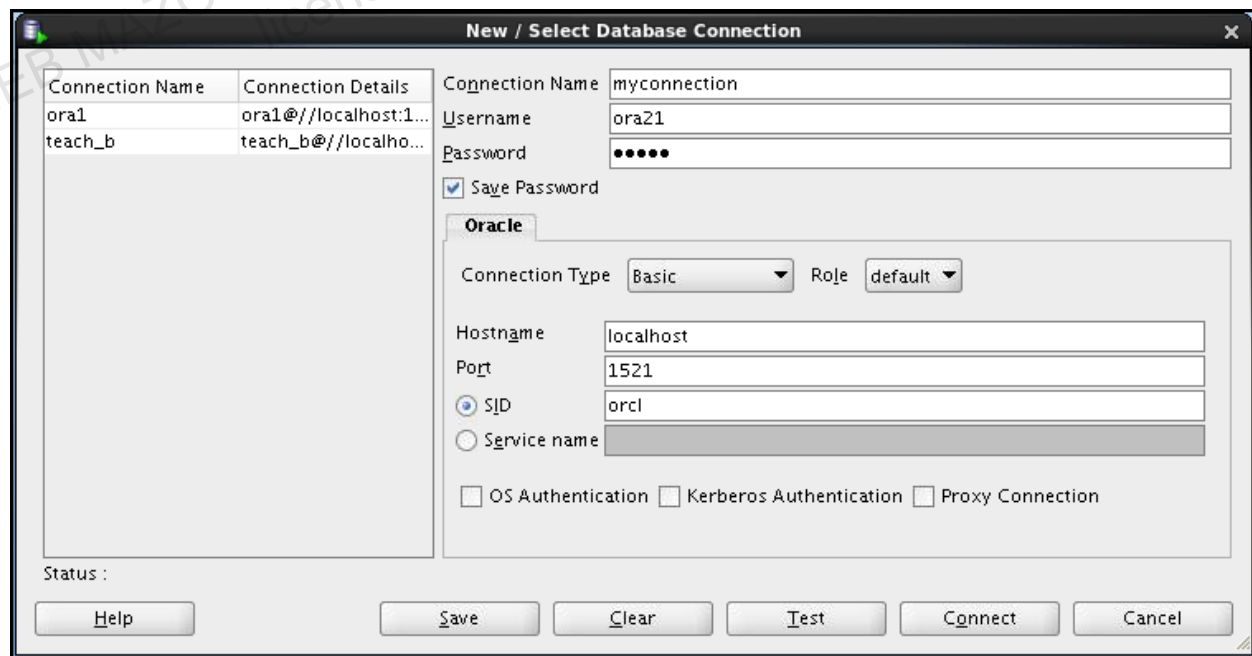
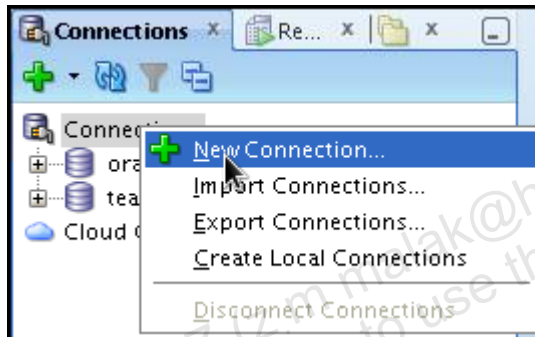
Solution 1-1: Using SQL Developer

1. Start SQL Developer by using the desktop icon.



2. Create a database connection using the following information:

- Connection Name: myconnection
- Username: ora21
- Password: ora21
- Hostname: localhost
- Port: 1521
- SID: orcl (or the value provided to you by the instructor)



3. Test the new connection. If the status is Success, connect to the database by using this new connection.
 - a. Click the Test button in the New/Select Database Connection window.

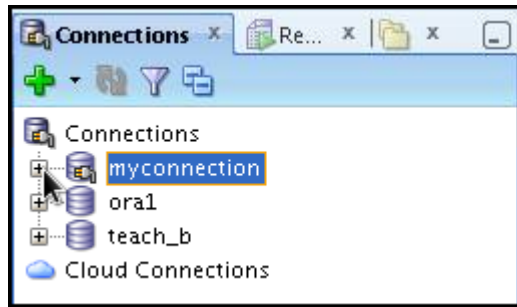
The screenshot shows the 'New / Select Database Connection' dialog box. On the left, there is a table with two columns: 'Connection Name' and 'Connection Details'. It lists 'ora1' and 'teach_b' with their respective connection details. The main area contains fields for 'Connection Name' (myconnection), 'Username' (ora21), and 'Password' (masked with dots). There is a checked 'Save Password' checkbox. Under the 'Oracle' tab, 'Connection Type' is set to 'Basic' and 'Role' to 'default'. The 'Hostname' is 'localhost', 'Port' is '1521', and 'SID' is 'orcl'. There are also checkboxes for 'OS Authentication', 'Kerberos Authentication', and 'Proxy Connection', all of which are unchecked. At the bottom, there are buttons for 'Help', 'Save', 'Clear', 'Test' (highlighted with a yellow border and a mouse cursor), 'Connect', and 'Cancel'. The status at the bottom left is 'Status : Success'.

- b. If the status is Success, click the Connect button.

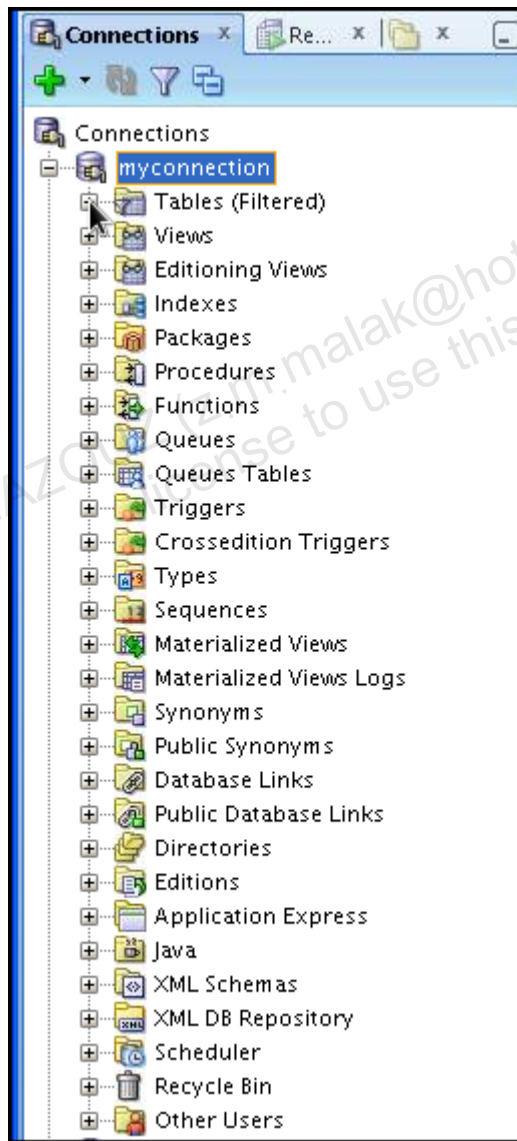
This screenshot is identical to the previous one, showing the 'New / Select Database Connection' dialog box with the same configuration. In this instance, the 'Connect' button is highlighted with a yellow border and a mouse cursor, indicating the next step in the process.

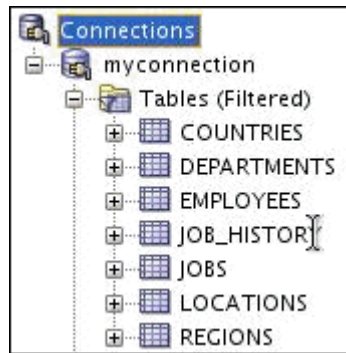
4. Browse the structure of the EMPLOYEES table and display its data.

- a. Expand the myconnection connection by clicking the plus sign next to it.



- b. Expand the Tables icon by clicking the plus sign next to it.





c. Display the structure of the **EMPLOYEES** table.

Click the **EMPLOYEES** table. The Columns tab displays the columns in the **EMPLOYEES** table as follows:

Columns						
	COLUMN_NAME	DATA_TYPE	NULLABLE	DATA_DEFAULT	COLUMN_ID	COM
1	EMPLOYEE_ID	NUMBER(6,0)	No	(null)	1	Primary
2	FIRST_NAME	VARCHAR2(20 BYTE)	Yes	(null)	2	First
3	LAST_NAME	VARCHAR2(25 BYTE)	No	(null)	3	Last n
4	EMAIL	VARCHAR2(25 BYTE)	No	(null)	4	Email
5	PHONE_NUMBER	VARCHAR2(20 BYTE)	Yes	(null)	5	Phone
6	HIRE_DATE	TIMESTAMP(7)	No	(null)	6	Date w
7	JOB_ID	VARCHAR2(10 BYTE)	No	(null)	7	Current
8	SALARY	NUMBER(8,2)	Yes	(null)	8	Monthl
9	COMMISSION_PCT	NUMBER(2,2)	Yes	(null)	9	Commis
10	MANAGER_ID	NUMBER(6,0)	Yes	(null)	10	Manage
11	DEPARTMENT_ID	NUMBER(4,0)	Yes	(null)	11	Depart

d. View the data of the **DEPARTMENTS** table.

In the Connections navigator, click the **DEPARTMENTS** table. Then click the Data tab.

	DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
1	10	Administration	200	1700
2	20	Marketing	201	1800
3	30	Purchasing	114	1700
4	40	Human Resources	203	2400
5	50	Shipping	121	1500
6	60	IT	103	1400
7	70	Public Relations	204	2700

...

5. Execute some basic `SELECT` statements to query the data in the `EMPLOYEES` table in the SQL Worksheet area. Use both the Execute Statement (and press F9) and the Run Script icons (or press F5) to execute the `SELECT` statements. Review the results of both methods of executing the `SELECT` statements on the appropriate tabbed pages.
 - a. Write a query to select the last name and salary for any employee whose salary is less than or equal to \$3,000.

```
SELECT last_name, salary
FROM employees
WHERE salary <= 3000;
```

	LAST_NAME	SALARY
1	Baida	2900
2	Tobias	2800
3	Himuro	2600
4	Colmenares	2500
5	Mikkilineni	2700
6	Landry	2400
7	Markle	2200
8	Atkinson	2800
9	Marlow	2500
10	Olson	2100
11	Rogers	2900
12	Gee	2400

...

- b. Write a query to display last name, job ID, and commission for all employees who are not entitled to receive a commission.

```
SELECT last_name, job_id, commission_pct
FROM employees
WHERE commission_pct IS NULL;
```

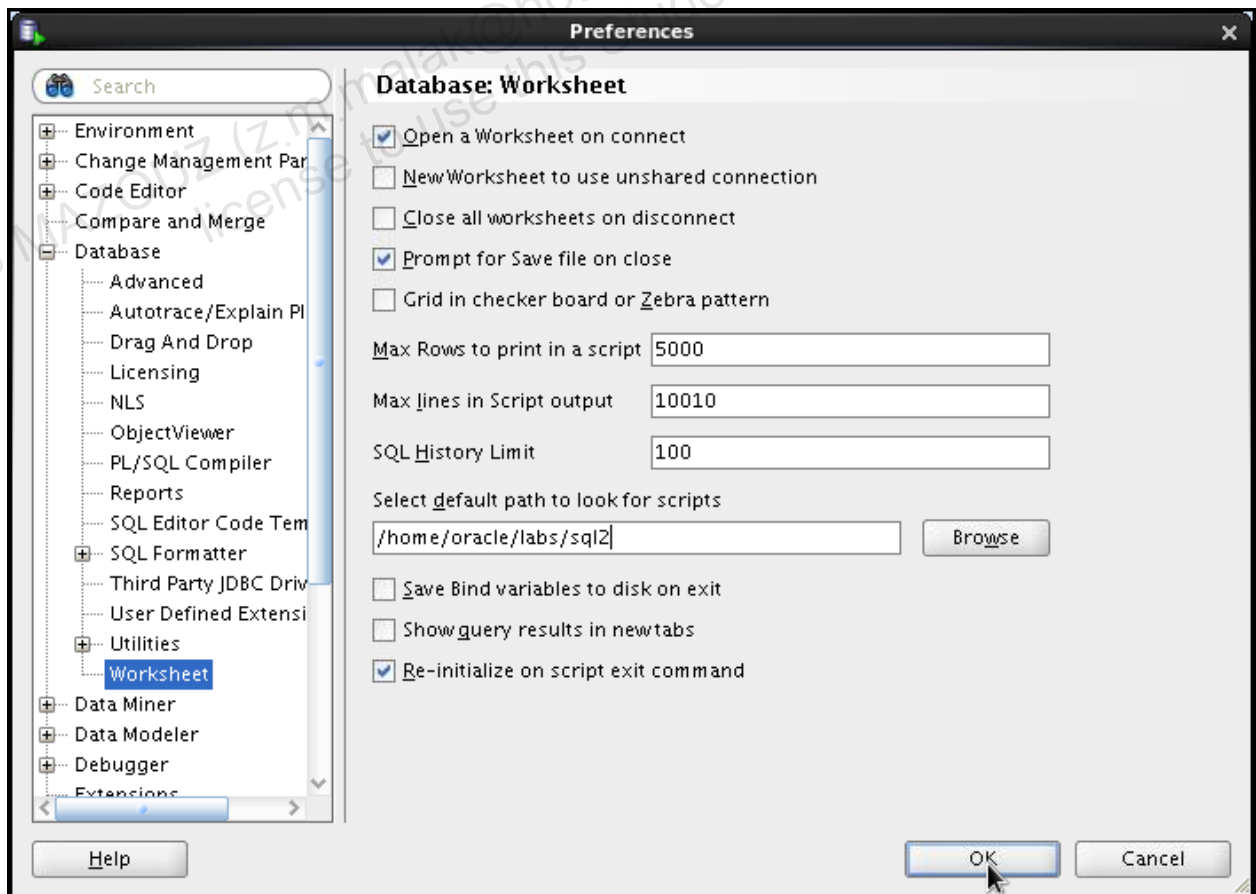
	LAST_NAME	JOB_ID	COMMISSION_PCT
1	King	AD_PRES	(null)
2	Kochhar	AD_VP	(null)
3	De Haan	AD_VP	(null)
4	Hunold	IT_PROG	(null)
5	Ernst	IT_PROG	(null)
6	Austin	IT_PROG	(null)
7	Pataballa	IT_PROG	(null)
8	Lorentz	IT_PROG	(null)
9	Greenberg	FI_MGR	(null)
10	Faviet	FI_ACCOUNT	(null)
11	Chen	FI_ACCOUNT	(null)

...

6. Set your script pathing preference to /home/oracle/labs/sql2.
 - a. Select Tools > Preferences > Database > Worksheet.
 - b. Enter the value in the "Select default path to look for scripts" field. Then, click OK.

Note: To view the number of rows selected, enable the feedback option and set it to 1.

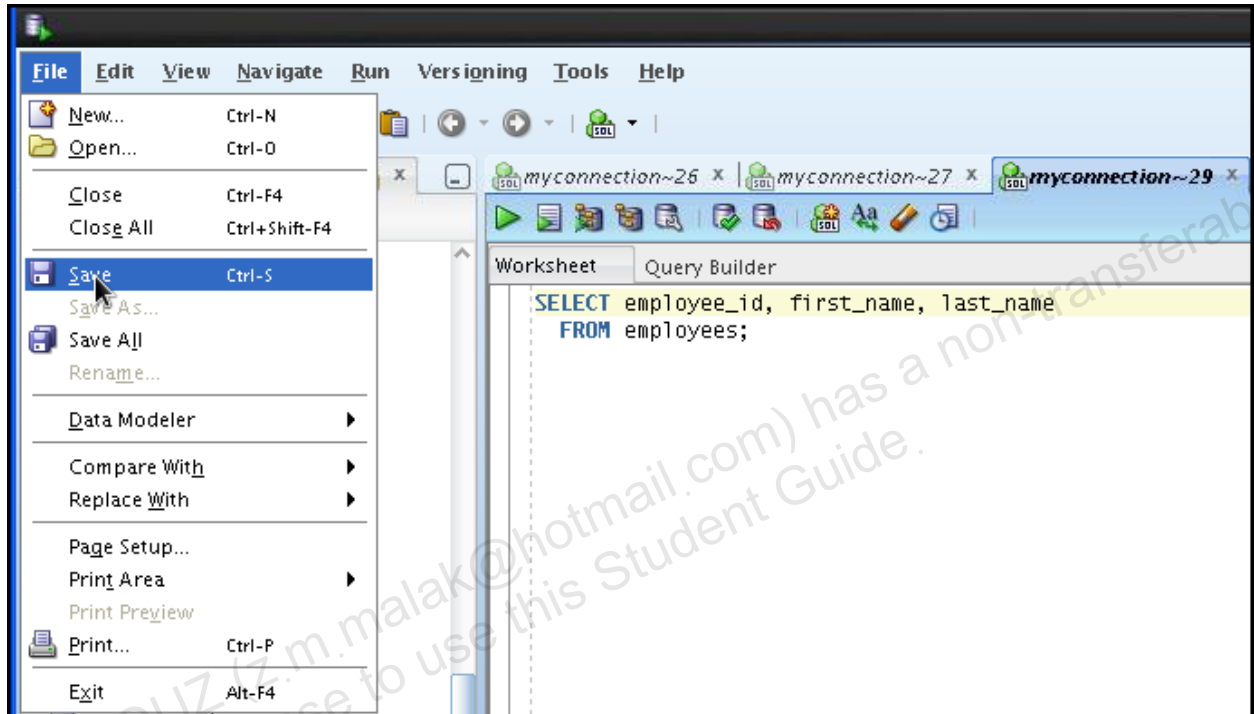
```
set feedback on;
set feedback 1;
```




7. Enter the following SQL statement:

```
SELECT employee_id, first_name, last_name
FROM employees;
```

8. Save the SQL statement to a script file by using the File > Save As menu item.
a. Select File > Save.




- b. Name the file `intro_test.sql`.
Enter `intro_test.sql` in the File_name text box.
c. Place the file under the `/home/oracle/labs/sql2/labs` folder.



Then, click Save.

9. Open and run `confidence.sql` from your `/home/oracle/labs/sql2/labs` folder and observe the output.

Open the confidence.sql script file by using the File > Open menu item.



Then, press F5 to execute the script.

The following is the expected result:

COUNT(*)

8
COUNT(*)

107
COUNT(*)

25
COUNT(*)

4

COUNT(*)

23
COUNT(*)

27
COUNT(*)

19
COUNT(*)

10

YASSINE MAZOUZ (z.m.malak@hotmail.com) has a non-transferable license to use this Student Guide.

Practices for Lesson 2: Introduction to Data Dictionary Views

Chapter 2

Practices for Lesson 2: Overview

Practice overview

This practice covers the following topics:

- Querying the dictionary views for table and column information
- Querying the dictionary views for constraint information
- Adding a comment to a table and querying the dictionary views for comment information

Practice 2-1: Introduction to Data Dictionary Views

Overview

In this practice, you query the dictionary views to find information about objects in your schema.

Tasks

1. Query the `USER_TABLES` data dictionary view to see information about the tables that you own.

	TABLE_NAME
1	REGIONS
2	LOCATIONS
3	DEPARTMENTS
4	JOBS
5	EMPLOYEES

...

2. Query the `ALL_TABLES` data dictionary view to see information about all the tables that you can access. Exclude the tables that you own.

Note: Your list may not exactly match the following list:

	TABLE_NAME	OWNER
1	DUAL	SYS
2	SYSTEM_PRIVILEGE_MAP	SYS
3	TABLE_PRIVILEGE_MAP	SYS
4	USER_PRIVILEGE_MAP	SYS
5	STMT_AUDIT_OPTION_MAP	SYS
6	AUDIT_ACTIONS	SYS
7	WRR\$_REPLAY_CALL_FILTER	SYS
8	HS_BULKLOAD_VIEW_OBJ	SYS
9	HS\$_PARALLEL_METADATA	SYS
10	HS_PARTITION_COL_NAME	SYS
11	HS_PARTITION_COL_TYPE	SYS

...

98	SDO_TOPO_DATA\$	MDSYS
99	SDO_GR_MOSAIC_0	MDSYS
100	SDO_GR_MOSAIC_1	MDSYS
101	SDO_GR_MOSAIC_2	MDSYS
102	SDO_GR_MOSAIC_3	MDSYS
103	SDO_GR_PARALLEL	MDSYS
104	SDO_GR_RDT_1	MDSYS
105	SDO_WFS_LOCAL_TXNS	MDSYS

- For a specified table, create a script that reports the column names, data types, and data types' lengths, as well as whether nulls are allowed. Prompt the user to enter the table name. Give appropriate aliases to the DATA_PRECISION and DATA_SCALE columns. Save this script in a file named lab_02_03.sql. For example, if the user enters DEPARTMENTS, the following output results:

	COLUMN_NAME	DATA_TYPE	DATA_LENGTH	PRECISION	SCALE	NULLABLE
1	DEPARTMENT_ID	NUMBER	22	4	0	N
2	DEPARTMENT_NAME	VARCHAR2	30	(null)	(null)	N
3	MANAGER_ID	NUMBER	22	6	0	Y
4	LOCATION_ID	NUMBER	22	4	0	Y

- Create a script that reports the column name, constraint name, constraint type, search condition, and status for a specified table. You must join the USER_CONSTRAINTS and USER_CONS_COLUMNS tables to obtain all this information. Prompt the user to enter the table name. Save the script in a file named lab_02_04.sql. For example, if the user enters DEPARTMENTS, the following output results:

	COLUMN_NAME	CONSTRAINT_NAME	CONSTRAINT_TYPE	SEARCH_CONDITION	STATUS
1	MANAGER_ID	DEPT_MGR_FK	R	(null)	ENABLED
2	LOCATION_ID	DEPT_LOC_FK	R	(null)	ENABLED
3	DEPARTMENT_ID	DEPT_ID_PK	P	(null)	ENABLED
4	DEPARTMENT_NAME	DEPT_NAME_NN	C	"DEPARTMENT_NAME" IS NOT NULL	ENABLED

- Add a comment to the DEPARTMENTS table. Then query the USER_TAB_COMMENTS view to verify that the comment is present.

COMMENTS
1 Company department information including name, code, and location.

- Run the lab_02_06_tab.sql script as a prerequisite for exercises 6 through 9. Alternatively, open the script file to copy the code and paste it into your SQL Worksheet. Then execute the script. This script:
 - Drops the existing DEPT2 and EMP2 tables
 - Creates the DEPT2 and EMP2 tables

Note: In Practice 2, you should have already dropped the DEPT2 and EMP2 tables so that they cannot be restored.

7. Confirm that both the DEPT2 and EMP2 tables are stored in the data dictionary.

	TABLE_NAME
1	DEPT2
2	EMP2

8. Confirm that the constraints were added, by querying the USER_CONSTRAINTS view. Note the types and names of the constraints.

	CONSTRAINT_NAME	CONSTRAINT_TYPE
1	MY_EMP_DEPT_ID_FK	R
2	MY_DEPT_ID_PK	P
3	MY_EMP_ID_PK	P

9. Display the object names and types from the USER_OBJECTS data dictionary view for the EMP2 and DEPT2 tables.

	OBJECT_NAME	OBJECT_TYPE
1	DEPT2	TABLE
2	EMP2	TABLE

Solution 2-1: Introduction to Data Dictionary Views

Solution

1. Query the data dictionary to see information about the tables you own.

```
SELECT table_name
FROM   user_tables;
```

2. Query the dictionary view to see information about all the tables that you can access. Exclude tables that you own.

```
SELECT table_name, owner
FROM   all_tables
WHERE  owner <> 'ORAXX';
```

3. For a specified table, create a script that reports the column names, data types, and data types' lengths, as well as whether nulls are allowed. Prompt the user to enter the table name. Give appropriate aliases to the DATA_PRECISION and DATA_SCALE columns. Save this script in a file named lab_02_03.sql.

```
SELECT column_name, data_type, data_length,
       data_precision PRECISION, data_scale SCALE, nullable
FROM   user_tab_columns
WHERE  table_name = UPPER('&tab_name');
```

To test, run the script and enter DEPARTMENTS as the table name.

4. Create a script that reports the column name, constraint name, constraint type, search condition, and status for a specified table. You must join the USER_CONSTRAINTS and USER_CONS_COLUMNS tables to obtain all this information. Prompt the user to enter the table name. Save the script in a file named lab_02_04.sql.

```
SELECT ucc.column_name, uc.constraint_name, uc.constraint_type,
       uc.search_condition, uc.status
FROM   user_constraints uc JOIN user_cons_columns ucc
ON     uc.table_name = ucc.table_name
AND    uc.constraint_name = ucc.constraint_name
AND    uc.table_name = UPPER('&tab_name');
```

To test, run the script and enter DEPARTMENTS as the table name.

5. Add a comment to the DEPARTMENTS table. Then query the USER_TAB_COMMENTS view to verify that the comment is present.

```
COMMENT ON TABLE departments IS
    'Company department information including name, code, and
    location.';

SELECT COMMENTS
FROM    user_tab_comments
WHERE   table_name = 'DEPARTMENTS';
```

6. Run the lab_02_06_tab.sql script as a prerequisite for exercises 6 through 9. Alternatively, open the script file to copy the code and paste it into your SQL Worksheet. Then execute the script. This script:

- Drops the DEPT2 and EMP2 tables
- Creates the DEPT2 and EMP2 tables

7. Confirm that both the DEPT2 and EMP2 tables are stored in the data dictionary.

```
SELECT    table_name
FROM      user_tables
WHERE     table_name IN ('DEPT2', 'EMP2');
```

8. Query the data dictionary to find out the constraint names and types for both the tables.

```
SELECT    constraint_name, constraint_type
FROM      user_constraints
WHERE     table_name IN ('EMP2', 'DEPT2');
```

9. Display the object names and types from the USER_OBJECTS data dictionary view for the EMP2 and DEPT2 tables.

```
SELECT    object_name, object_type
FROM      user_objects
WHERE     object_name= 'EMP2'
OR        object_name= 'DEPT2';
```

YASSINE MAZOUZ (z.m.malak@hotmail.com) has a non-transferable license to use this Student Guide.

Practices for Lesson 3: Creating Sequences, Synonyms, and Indexes

Chapter 3

Practices for Lesson 3: Overview

Practices Overview

This practice covers the following topics:

- Creating sequences
- Using sequences
- Querying the dictionary views for sequence information
- Creating synonyms
- Querying the dictionary views for synonyms information
- Creating indexes
- Querying the dictionary views for indexes information

Note: Before starting this practice, execute
/home/oracle/sql2/code_ex/code_ex_scripts/clean_up_scripts/cleanup_03.sql script.

Unauthorized reproduction or distribution prohibited. Copyright© 2016, Oracle and/or its affiliates.

Tasks

- | Column Name | ID | NAME |
|-------------|-------------|----------|
| Key Type | Primary key | |
| Null/Unique | | |
| FK Table | | |
| FK Column | | |
| Data Type | NUMBER | VARCHAR2 |
| Length | 7 | 25 |

- [illegible]

- | | SYNONYM_NAME | TABLE_OWNER | TABLE_NAME | DB_LINK |
|---|--------------|-------------|------------|---------|
| 1 | EMP1 | ORA21 | EMPLOYEES | (null) |

- Chapter 3 - Page 3

8. Create the `SALES_DEPT` table based on the following table instance chart. Name the index for the `PRIMARY KEY` column `SALES_PK_IDX`. Then query the data dictionary view to find the index name, table name, and whether the index is unique.

Column Name	Team_Id	Location
Primary Key	Yes	
Data Type	Number	VARCHAR2
Length	3	30

INDEX_NAME	TABLE_NAME	UNIQUENESS
1 SALES_PK_IDX	SALES_DEPT	NONUNIQUE

9. Drop the tables and sequences created in this practice.

Solution 3-1: Creating Sequences, Synonyms, and Indexes

1. Create the DEPT table based on the following table instance chart. Confirm that the table is created.

Column Name	ID	NAME
Key Type	Primary key	
Null/Unique		
FK Table		
FK Column		
Data Type	NUMBER	VARCHAR2
Length	7	25

```
CREATE TABLE dept
(id NUMBER(7) CONSTRAINT department_id_pk PRIMARY KEY,
name VARCHAR2(25));
```

To confirm that the table was created and to view its structure, issue the following command:

```
DESCRIBE dept;
```

2. You need a sequence that can be used with the primary key column of the DEPT table. The sequence should start at 200 and have a maximum value of 1,000. Have your sequence increment by 10. Name the sequence DEPT_ID_SEQ.

```
CREATE SEQUENCE dept_id_seq
START WITH 200
INCREMENT BY 10
MAXVALUE 1000;
```

3. To test your sequence, write a script to insert two rows in the DEPT table. Name your script lab_03_03.sql. Be sure to use the sequence that you created for the ID column. Add two departments: Education and Administration. Confirm your additions. Run the commands in your script.

```
INSERT INTO dept
VALUES (dept_id_seq.nextval, 'Education');
INSERT INTO dept
VALUES (dept_id_seq.nextval, 'Administration');
```

- Find the names of your sequences. Write a query in a script to display the following information about your sequences: sequence name, maximum value, increment size, and last number. Name the script lab_03_04.sql. Run the statement in your script.

```
SELECT  sequence_name, max_value, increment_by, last_number
FROM    user_sequences;
```

- Create a synonym for your EMPLOYEES table. Call it EMP1. Then find the names of all synonyms that are in your schema.

```
CREATE SYNONYM emp1 FOR EMPLOYEES;
SELECT *
FROM    user_synonyms;
```

- Drop the EMP1 synonym.

```
DROP SYNONYM emp1;
```

- Create a nonunique index on the NAME column in the DEPT table.

```
CREATE INDEX dept_name_idx ON dept (name);
```

- Create the SALES_DEPT table based on the following table instance chart. Name the index for the PRIMARY KEY column SALES_PK_IDX. Then query the data dictionary view to find the index name, table name, and whether the index is unique.

Column Name	Team_Id	Location
Primary Key	Yes	
Data Type	Number	VARCHAR2
Length	3	30

```
CREATE TABLE SALES_DEPT
(team_id NUMBER(3)
PRIMARY KEY USING INDEX
(CREATE INDEX sales_pk_idx ON
SALES_DEPT(team_id)),
location VARCHAR2(30));
SELECT INDEX_NAME, TABLE_NAME, UNIQUENESS
FROM USER_INDEXES
WHERE TABLE_NAME = 'SALES_DEPT';
```

- Drop the tables and sequences created in this practice.

```
DROP TABLE DEPT;
DROP TABLE SALES_DEPT;
DROP SEQUENCE dept_id_seq;
```


Practices for Lesson 4: Creating Views

Chapter 4

Practices for Lesson 4: Overview

Practices Overview

This practice covers the following topics:

- Creating a simple view
- Creating a complex view
- Creating a view with a check constraint
- Attempting to modify data in the view
- Querying the dictionary views for view information
- Removing views

Practice 4-1: Creating Views

Overview:

This lesson's practice provides you with a variety of exercises in creating, using, querying data dictionary views for view information, and removing views.

Tasks:

1. The staff in the HR department wants to hide some of the data in the `EMPLOYEES` table. Create a view called `EMPLOYEES_VU` based on the employee numbers, employee last names, and department numbers from the `EMPLOYEES` table. The heading for the employee name should be `EMPLOYEE`.
2. Confirm that the view works. Display the contents of the `EMPLOYEES_VU` view.

	EMPLOYEE_ID	EMPLOYEE	DEPARTMENT_ID
1	100	King	90
2	101	Kochhar	90
3	102	De Haan	90
4	103	Hunold	60
5	104	Ernst	60
6	105	Austin	60
7	106	Pataballa	60
8	107	Lorentz	60
9	108	Greenberg	100
10	109	Faviet	100
11	110	Chen	100
12	111	Sciarra	100
13	112	Urman	100

...

3. Using your `EMPLOYEES_VU` view, write a query for the HR department to display all employee names and department numbers.

	EMPLOYEE	DEPARTMENT_ID
1	King	90
2	Kochhar	90
3	De Haan	90
4	Hunold	60
5	Ernst	60
6	Austin	60
7	Pataballa	60
8	Lorentz	60
9	Greenberg	100
10	Faviet	100
11	Chen	100

...

4. Department 80 needs access to its employee data. Create a view named `DEPT50` that contains the employee numbers, employee last names, and department numbers for all employees in department 80. You have been asked to label the view columns `EMPNO`, `EMPLOYEE`, and `DEPTNO`. For security purposes, do not allow an employee to be reassigned to another department through the view.
5. Display the structure and contents of the `DEPT80` view.

```
DESCRIBE dept80
Name      Null    Type
-----
EMPNO     NOT NULL NUMBER(6)
EMPLOYEE  NOT NULL VARCHAR2(25)
DEPTNO                    NUMBER(4)
```

	EMPNO	EMPLOYEE	DEPTNO
1	145	Russell	80
2	146	Partners	80
3	147	Errazuriz	80
4	148	Cambrault	80
5	149	Zlotkey	80
6	150	Tucker	80
7	151	Bernstein	80
8	152	Hall	80
9	153	Olsen	80
10	154	Cambrault	80
11	155	Tuvault	80

...

6. Test your view. Attempt to reassign Abel to department 80.

```
Error report:
SQL Error: ORA-01402: view WITH CHECK OPTION where-clause violation
01402. 00000 - "view WITH CHECK OPTION where-clause violation"
*Cause:
*Action:
```

7. Run `lab_04_07.sql` to create the `dept50` view for this exercise. You need to determine the names and definitions of all the views in your schema. Create a report that retrieves view information: the view name and text from the `USER_VIEWS` data dictionary view.
Note: The `EMP_DETAILS_VIEW` was created as part of your schema.
Note: You can see the complete definition of the view if you use Run Script (or press F5) in SQL Developer. If you use Execute Statement (or press F9) in SQL Developer, scroll horizontally in the result pane. If you use SQL*Plus, to see more contents of a `LONG` column, use the `SET LONG n` command, where `n` is the value of the number of characters of the `LONG` column that you want to see.

VIEW_NAME	TEXT
1 DEPT50	SELECT employee_id empno, last_name employee, department_id deptno FROM
2 DEPT80	SELECT employee_id empno, last_name employee, department_id deptno FROM
3 EMPLOYEES_VU	SELECT employee_id, last_name employee, department_id FROM employees
4 EMP_DETAILS_VIEW	SELECT e.employee_id, e.job_id, e.manager_id, e.department_id, d.location_id, l.count

8. Remove the views created in this practice.

Solution 4-1: Creating Views

1. The staff in the HR department wants to hide some of the data in the `EMPLOYEES` table. Create a view called `EMPLOYEES_VU` based on the employee numbers, employee last names, and department numbers from the `EMPLOYEES` table. The heading for the employee name should be `EMPLOYEE`.

```
CREATE OR REPLACE VIEW employees_vu AS
    SELECT employee_id, last_name employee, department_id
    FROM employees;
```

2. Confirm that the view works. Display the contents of the `EMPLOYEES_VU` view.

```
SELECT *
FROM employees_vu;
```

3. Using your `EMPLOYEES_VU` view, write a query for the HR department to display all employee names and department numbers.

```
SELECT employee, department_id
FROM employees_vu;
```

4. Department 80 needs access to its employee data. Create a view named `DEPT80` that contains the employee numbers, employee last names, and department numbers for all employees in department 80. They have requested that you label the view columns `EMPNO`, `EMPLOYEE`, and `DEPTNO`. For security purposes, do not allow an employee to be reassigned to another department through the view.

```
CREATE VIEW dept80 AS
    SELECT employee_id empno, last_name employee,
           department_id deptno
    FROM employees
    WHERE department_id = 80
    WITH CHECK OPTION CONSTRAINT emp_dept_80;
```

5. Display the structure and contents of the `DEPT80` view.

```
DESCRIBE dept80

SELECT *
FROM dept80;
```

6. Test your view. Attempt to reassign Abel to department 50.

```
UPDATE    dept80
SET       deptno = 50
WHERE     employee = 'Abel';
```

The error is because the DEPT50 view has been created with the WITH CHECK OPTION constraint. This ensures that the DEPTNO column in the view is protected from being changed.

7. Run lab_04_07.sql to create the dept50 view for this exercise. You need to determine the names and definitions of all the views in your schema. Create a report that retrieves view information: the view name and text from the USER_VIEWS data dictionary view.

Note: The EMP_DETAILS_VIEW was created as part of your schema.

Note: You can see the complete definition of the view if you use Run Script (or press F5) in SQL Developer. If you use Execute Statement (or press F9) in SQL Developer, scroll horizontally in the result pane. If you use SQL*Plus to see more contents of a LONG column, use the SET LONG *n* command, where *n* is the value of the number of characters of the LONG column that you want to see.

```
SELECT    view_name, text
FROM      user_views;
```

8. Remove the views created in this practice.

```
DROP VIEW employees_vu;
DROP VIEW dept80;
DROP VIEW dept50;
```

Youssef MAZOUZ (z.m.malak@hotmail.com) has a non-transferable
license to use this Student Guide.

Practices for Lesson 5: Managing Schema Objects

Chapter 5

Practices for Lesson 5: Overview

Practice Overview

This practice covers the following topics:

- Adding and dropping constraints
- Deferring constraints
- Creating external tables

Note: Before starting this practice, execute `/home/oracle/labs/sql2/code_ex/cleanup_scripts/cleanup_05.sql` script.

Practice 5-1: Managing Schema Objects

Overview

In this practice, you add, drop, and defer constraints. You create external tables.

Note: Execute `cleanup_05.sql` script from `/home/oracle/labs/sql2/code_ex/cleanup_scripts/` before performing the following tasks.

Tasks

1. Create the `DEPT2` table based on the following table instance chart. Enter the syntax in the SQL Worksheet. Then, execute the statement to create the table. Confirm that the table is created.

Column Name	ID	NAME
Key Type		
Nulls/Unique		
FK Table		
FK Column		
Data type	NUMBER	VARCHAR2
Length	7	25

```
DESCRIBE dept2
Name Null Type
-----
ID      NUMBER(7)
NAME    VARCHAR2(25)
```

- Populate the DEPT2 table with data from the DEPARTMENTS table. Include only the columns that you need. Confirm that the rows are inserted.

R2	ID	R2	NAME
1	10		Administration
2	20		Marketing
3	30		Purchasing
4	40		Human Resources
5	50		Shipping
6	60		IT
7	70		Public Relations
8	80		Sales
9	90		Executive
10	100		Finance
11	110		Accounting
12	120		Treasury
13	130		Corporate Tax
14	140		Control And Credit
15	150		Shareholder Services

...

- Create the EMP2 table based on the following table instance chart. Enter the syntax in the SQL Worksheet. Then execute the statement to create the table. Confirm that the table is created.

Column Name	ID	LAST_NAME	FIRST_NAME	DEPT_ID
Key Type				
Nulls/Unique				
FK Table				
FK Column				
Data type	NUMBER	VARCHAR2	VARCHAR2	NUMBER
Length	7	25	25	7

DESCRIBE emp2	
Name	Null Type

ID	NUMBER(7)
LAST_NAME	VARCHAR2(25)
FIRST_NAME	VARCHAR2(25)
DEPT_ID	NUMBER(7)

4. Add a table-level PRIMARY KEY constraint to the EMP2 table on the ID column. The constraint should be named at creation. Name the constraint my_emp_id_pk.
5. Create a PRIMARY KEY constraint to the DEPT2 table using the ID column. The constraint should be named at creation. Name the constraint my_dept_id_pk.
6. Add a foreign key reference on the EMP2 table that ensures that the employee is not assigned to a nonexistent department. Name the constraint my_emp_dept_id_fk.
7. Modify the EMP2 table. Add a COMMISSION column of the NUMBER data type, precision 2, scale 2. Add a constraint to the COMMISSION column that ensures that a commission value is greater than zero.
8. Drop the EMP2 and DEPT2 tables so that they cannot be restored.
9. Create an external table library_items_ext. Use the ORACLE_LOADER access driver.

Note: The emp_dir directory and library_items.dat file are already created for this exercise. library_items.dat has records in the following format:

```
2354,      2264, 13.21, 150,
2355,      2289, 46.23, 200,
2355,      2264, 50.00, 100,
```

- a. Open the lab_05_09.sql file. Observe the code snippet to create the library_items_ext external table. Then replace <TODO1>, <TODO2>, <TODO3>, and <TODO4> as appropriate and save the file as lab_05_09_soln.sql. Run the script to create the external table.
- b. Query the library_items_ext table.

	Q1	Q2	Q3	Q4	Q5
		CATEGORY_ID	BOOK_ID	BOOK_PRICE	QUANTITY
1		2354	2264	13.21	150
2		2355	2289	46.23	200
3		2355	2264	50	100

10. The HR department needs a report of the addresses of all departments. Create an external table as dept_add_ext using the ORACLE_DATAPUMP access driver. The report should show the location ID, street address, city, state or province, and country in the output. Use a NATURAL JOIN to produce the results.

Note: The emp_dir directory is already created for this exercise.

- a. Open the lab_05_10.sql file. Observe the code snippet to create the dept_add_ext external table. Then, replace <TODO1>, <TODO2>, and <TODO3> with the appropriate code. Replace <oraxx_emp4.exp> and <oraxx_emp5.exp> with the appropriate file names. For example, if you are the ora21 user, your file names are ora21_emp4.exp and ora21_emp5.exp. Save the script as lab_05_10_soln.sql.
- b. Run the lab_05_10_soln.sql script to create the external table.
- c. Query the dept_add_ext table.

LOCATION_ID	STREET_ADDRESS	CITY	STATE_PROVINCE	COUNTRY_NAME
1	1000 1297 Via Cola di Rie	Roma	(null)	Italy
2	1100 93091 Calle della Testa	Venice	(null)	Italy
3	1200 2017 Shinjuku-ku	Tokyo	Tokyo Prefecture	Japan
4	1300 9450 Kamiya-cho	Hiroshima	(null)	Japan
5	1400 2014 Jabberwocky Rd	Southlake	Texas	United States of America
6	1500 2011 Interiors Blvd	South San Francisco	California	United States of America
7	1600 2007 Zagora St	South Brunswick	New Jersey	United States of America
8	1700 2004 Charade Rd	Seattle	Washington	United States of America
9	1800 147 Spadina Ave	Toronto	Ontario	Canada
10	1900 6092 Boxwood St	Whitehorse	Yukon	Canada

Note: When you perform the preceding step, two files `oraxx_emp4.exp` and `oraxx_emp5.exp` are created under the default directory `emp_dir`.

11. Create the `emp_books` table and populate it with data. Set the primary key as deferred and observe what happens at the end of the transaction.

- a. Run the `lab_05_11_a.sql` file to create the `emp_books` table. Observe that the `emp_books_pk` primary key is not created as deferrable.

```
table EMP_BOOKS created.
```

- b. Run the `lab_05_11_b.sql` file to populate data into the `emp_books` table. What do you observe?

```
1 rows inserted.

Error starting at line 2 in command:
insert into emp_books values(300,'Change Management')
Error report:
SQL Error: ORA-00001: unique constraint (ORA21.EMP_BOOKS_PK) violated
00001. 00000 - "unique constraint (%s.%s) violated"
*Cause:   An UPDATE or INSERT statement attempted to insert a duplicate key.
          For Trusted Oracle configured in DBMS MAC mode, you may see
          this message if a duplicate entry exists at a different level.
*Action:  Either remove the unique restriction or do not insert the key.
```

- c. Set the `emp_books_pk` constraint as deferred. What do you observe?

```
Error starting at line 1 in command:
set constraint emp_books_pk deferred
Error report:
SQL Error: ORA-02447: cannot defer a constraint that is not deferrable
02447. 00000 - "cannot defer a constraint that is not deferrable"
*Cause:   An attempt was made to defer a nondeferrable constraint
*Action:  Drop the constraint and create a new one that is deferrable
```

- d. Drop the `emp_books_pk` constraint.

```
table EMP_BOOKS altered.
```

- e. Modify the `emp_books` table definition to add the `emp_books_pk` constraint as deferrable this time.

```
table EMP_BOOKS altered.
```

- f. Set the `emp_books_pk` constraint as deferred.

```
constraint EMP_BOOKS_PK succeeded.
```

- g. Run the `lab_05_11_g.sql` file to populate data into the `emp_books` table. What do you observe?

```
1 rows inserted
1 rows inserted
1 rows inserted
```

- h. Commit the transaction. What do you observe?

```
Error starting at line 1 in command:
commit
Error report:
SQL Error: ORA-02091: transaction rolled back
ORA-00001: unique constraint (ORA21.EMP_BOOKS_PK) violated
02091. 00000 - "transaction rolled back"
*Cause:      Also see error 2092. If the transaction is aborted at a remote
              site then you will only see 2091; if aborted at host then you will
              see 2092 and 2091.
*Action:     Add rollback segment and retry the transaction.
```

Solution 5-1: Managing Schema Objects

Solution

1. Create the DEPT2 table based on the following table instance chart. Enter the syntax in the SQL Worksheet. Then, execute the statement to create the table. Confirm that the table is created.

Column Name	ID	NAME
Key Type		
Nulls/Unique		
FK Table		
FK Column		
Data type	NUMBER	VARCHAR2
Length	7	25

```
CREATE TABLE dept2
  (id NUMBER(7),
   name VARCHAR2(25));

DESCRIBE dept2
```

2. Populate the DEPT2 table with data from the DEPARTMENTS table. Include only the columns that you need.

```
INSERT INTO dept2
SELECT department_id, department_name
FROM departments;
```

3. Create the EMP2 table based on the following table instance chart. Enter the syntax in the SQL Worksheet. Then execute the statement to create the table. Confirm that the table is created.

Column Name	ID	LAST_NAME	FIRST_NAME	DEPT_ID
Key Type				
Nulls/Unique				
FK Table				
FK Column				
Data type	NUMBER	VARCHAR2	VARCHAR2	NUMBER
Length	7	25	25	7


```
CREATE TABLE emp2
(id          NUMBER(7),
last_name    VARCHAR2(25),
first_name   VARCHAR2(25),
dept_id      NUMBER(7));

DESCRIBE emp2
```

4. Add a table-level PRIMARY KEY constraint to the EMP2 table on the ID column. The constraint should be named at creation. Name the constraint my_emp_id_pk.

```
ALTER TABLE emp2
ADD CONSTRAINT my_emp_id_pk PRIMARY KEY (id);
```

5. Create a PRIMARY KEY constraint to the DEPT2 table using the ID column. The constraint should be named at creation. Name the constraint my_dept_id_pk.

```
ALTER TABLE dept2
ADD CONSTRAINT my_dept_id_pk PRIMARY KEY(id);
```

6. Add a foreign key reference on the EMP2 table that ensures that the employee is not assigned to a nonexistent department. Name the constraint my_emp_dept_id_fk.

```
ALTER TABLE emp2
ADD CONSTRAINT my_emp_dept_id_fk
FOREIGN KEY (dept_id) REFERENCES dept2(id);
```

7. Modify the EMP2 table. Add a COMMISSION column of the NUMBER data type, precision 2, scale 2. Add a constraint to the COMMISSION column that ensures that a commission value is greater than zero.

```
ALTER TABLE emp2
ADD commission NUMBER(2,2)
CONSTRAINT my_emp_comm_ck CHECK (commission > 0);
```

8. Drop the EMP2 and DEPT2 tables so that they cannot be restored.

```
DROP TABLE emp2 PURGE;
DROP TABLE dept2 PURGE;
```

9. Create an external table `library_items_ext`. Use the `ORACLE_LOADER` access driver.

Note: The `emp_dir` directory and `library_items.dat` are already created for this exercise. Ensure that the external file and the database are on the same machine.

`library_items.dat` has records in the following format:

```
2354,      2264, 13.21, 150,
2355,      2289, 46.23, 200,
2355,      2264, 50.00, 100,
```

- a. Open the `lab_05_09.sql` file. Observe the code snippet to create the `library_items_ext` external table. Then, replace `<TODO1>`, `<TODO2>`, `<TODO3>`, and `<TODO4>` as appropriate and save the file as `lab_05_09_soln.sql`. Run the script to create the external table.

```
CREATE TABLE library_items_ext ( category_id    number(12)
                                , book_id number(6)
                                , book_price number(8,2)
                                , quantity    number(8)
                                )

ORGANIZATION EXTERNAL
  (TYPE ORACLE_LOADER
   DEFAULT DIRECTORY emp_dir
   ACCESS PARAMETERS (RECORDS DELIMITED BY NEWLINE
                     FIELDS TERMINATED BY ',')
   LOCATION ('library_items.dat')
  )
REJECT LIMIT UNLIMITED;
```

- b. Query the `library_items_ext` table.

```
SELECT * FROM library_items_ext;
```

10. The HR department needs a report of addresses of all the departments. Create an external table as `dept_add_ext` using the `ORACLE_DATAPUMP` access driver. The report should show the location ID, street address, city, state or province, and country in the output. Use a `NATURAL JOIN` to produce the results.

Note: The `emp_dir` directory is already created for this exercise. Ensure that the external file and the database are on the same machine.

- a. Open the `lab_05_10.sql` file. Observe the code snippet to create the `dept_add_ext` external table. Then, replace `<TODO1>`, `<TODO2>`, and `<TODO3>` with appropriate code. Replace `<oraxx_emp4.exp>` and `<oraxx_emp5.exp>` with appropriate file names. For example, if you are user `ora21`, your file names are `ora21_emp4.exp` and `ora21_emp5.exp`. Save the script as `lab_5_10_soln.sql`.

```
CREATE TABLE dept_add_ext (location_id,
                           street_address, city,
                           state_province,
                           country_name)

ORGANIZATION EXTERNAL(
  TYPE ORACLE_DATAPUMP
  DEFAULT DIRECTORY emp_dir
  LOCATION ('oraxx_emp4.exp','oraxx_emp5.exp'))
PARALLEL
AS
SELECT location_id, street_address, city, state_province,
country_name
FROM locations
NATURAL JOIN countries;
```

Note: When you perform the preceding step, two files `oraxx_emp4.exp` and `oraxx_emp5.exp` are created under the default directory `emp_dir`.

- b. Run the `lab_05_10_soln.sql` script to create the external table.
- c. Query the `dept_add_ext` table.

```
SELECT * FROM dept_add_ext;
```

11. Create the `emp_books` table and populate it with data. Set the primary key as deferred and observe what happens at the end of the transaction.

- a. Run the `lab_05_11_a.sql` script to create the `emp_books` table. Observe that the `emp_books_pk` primary key is not created as deferrable.

```
DROP TABLE emp_books CASCADE CONSTRAINTS;
CREATE TABLE emp_books (book_id number,
                          title varchar2(20), CONSTRAINT
emp_books_pk PRIMARY KEY (book_id));
```

- b. Run the `lab_05_11_b.sql` script to populate data into the `emp_books` table. What do you observe?

```
INSERT INTO emp_books VALUES(300,'Organizations');
INSERT INTO emp_books VALUES(300,'Change Management');
```

The first row is inserted. However, you see the `ora-00001` error with the second row insertion.

- c. Set the `emp_books_pk` constraint as deferred. What do you observe?

```
SET CONSTRAINT emp_books_pk DEFERRED;
```

You see the following error: “ORA-02447: Cannot defer a constraint that is not deferrable.”

- d. Drop the emp_books_pk constraint.

```
ALTER TABLE emp_books DROP CONSTRAINT emp_books_pk;
```

- e. Modify the emp_books table definition to add the emp_books_pk constraint as deferrable this time.

```
ALTER TABLE emp_books ADD (CONSTRAINT emp_books_pk PRIMARY KEY  
(book_id) DEFERRABLE);
```

- f. Set the emp_books_pk constraint as deferred.

```
SET CONSTRAINT emp_books_pk DEFERRED;
```

- g. Run the lab_05_11_g.sql script to populate data into the emp_books table.
What do you observe?

```
INSERT INTO emp_books VALUES (300, 'Change Management');  
INSERT INTO emp_books VALUES (300, 'Personality');  
INSERT INTO emp_books VALUES (350, 'Creativity');
```

You see that all the rows are inserted.

- h. Commit the transaction. What do you observe?

```
COMMIT;
```

You see that the transaction is rolled back by the database at this point, because the COMMIT failed due to the constraint violation.

Practices for Lesson 6: Retrieving Data by Using Subqueries

Chapter 6

Practices for Lesson 6: Overview

Practice Overview

This practice covers the following topics:

- Creating multiple-column subqueries
- Writing correlated subqueries
- Using the `EXISTS` operator
- Using scalar subqueries
- Using the `WITH` clause

Practice 6-1: Retrieving Data by Using Subqueries

Overview

In this practice, you write multiple-column subqueries, and correlated and scalar subqueries. You also solve problems by writing the `WITH` clause.

Tasks

1. Write a query to display the last name, department number, and salary of any employee whose department number and salary both match the department number and salary of any employee who earns a commission.

	LAST_NAME	DEPARTMENT_ID	SALARY
1	Russell	80	14000
2	Partners	80	13500
3	Errazuriz	80	12000
4	Abel	80	11000
5	Cambrault	80	11000
6	Vishney	80	10500
7	Zlotkey	80	10500
8	Bloom	80	10000
9	King	80	10000
10	Tucker	80	10000
11	Greene	80	9500

...

2. Display the last name, department name, and salary of any employee whose salary and `job_ID` match the salary and `job_ID` of any employee located in location ID 1700.

	LAST_NAME	DEPARTMENT_NAME	SALARY
1	Whalen	Administration	4400
2	Colmenares	Purchasing	2500
3	Himuro	Purchasing	2600
4	Tobias	Purchasing	2800
5	Baida	Purchasing	2900
6	Khoo	Purchasing	3100
7	Raphaely	Purchasing	11000
8	Grant	Shipping	2600
9	OConnell	Shipping	2600
10	Walsh	Shipping	3100
11	Jones	Shipping	2800

...

3. Create a query to display the last name, hire date, and salary for all employees who have the same salary and `manager_ID` as Kochhar.

Note: Do not display Kochhar in the result set.

	LAST_NAME	HIRE_DATE	SALARY
1	De Haan	13-JAN-01	17000

4. Create a query to display the employees who earn a salary that is higher than the salary of all the sales managers (`JOB_ID = 'SA_MAN'`). Sort the results from the highest to the lowest.

	LAST_NAME	JOB_ID	SALARY
1	King	AD_PRES	24000
2	De Haan	AD_VP	17000
3	Kochhar	AD_VP	17000

5. Display details such as the employee ID, last name, and department ID of those employees who live in cities the names of which begin with *T*.

	EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID
1	202	Fay	20
2	201	Hartstein	20

6. Write a query to find all employees who earn more than the average salary in their departments. Display last name, salary, department ID, and the average salary for the department. Sort by average salary and round to two decimals. Use aliases for the columns retrieved by the query as shown in the sample output.

	ENAME	SALARY	DEPTNO	DEPT_AVG
1	Fripp	8200	50	3475.56
2	Chung	3800	50	3475.56
3	Kaufling	7900	50	3475.56
4	Mourgos	5800	50	3475.56
5	Bell	4000	50	3475.56
6	Rajs	3500	50	3475.56
7	Everett	3900	50	3475.56
8	Sarchand	4200	50	3475.56
9	Bull	4100	50	3475.56
10	Vollman	6500	50	3475.56
11	Ladwig	3600	50	3475.56
12	Dilly	3600	50	3475.56
13	Weiss	8000	50	3475.56

...

7. Find all employees who are not supervisors.
- a. First, do this by using the NOT EXISTS operator.

	LAST_NAME
1	Abel
2	Ande
3	Atkinson
4	Austin
5	Baer
6	Baida
7	Banda
8	Bates
9	Bell
10	Bernstein
11	Bissot
12	Bloom
13	Bull
14	Cabrio
15	Cambrault

...

- b. Can this be done by using the NOT IN operator? How, or why not? If not, try out using another solution.

	LAST_NAME
1	Abel
2	Ande
3	Atkinson
4	Austin
5	Baer
6	Baida
7	Banda
8	Bates
9	Bell
10	Bernstein
11	Bissot
12	Bloom

...

8. Write a query to display the last names of the employees who earn less than the average salary in their departments.

	LAST_NAME
1	Chen
2	Sciarra
3	Urman
4	Popp
5	Khoo
6	Baida
7	Tobias
8	Himuro
9	Colmenares
10	Kochhar
11	De Haan
12	Fay
13	Gietz
14	Nayer

...

9. Write a query to display the last names of the employees who have one or more coworkers in their departments with later hire dates but higher salaries.

	LAST_NAME
1	De Haan
2	Austin
3	Lorentz
4	Pataballa
5	Faviet
6	Sciarra
7	Tobias
8	Bell
9	Sarchand
10	Rajs
11	Ladwig
12	Mallin
13	Kaufling

...

10. Write a query to display the employee ID, last names, and department names of all the employees.

Note: Use a scalar subquery to retrieve the department name in the `SELECT` statement.

	EMPLOYEE_ID	LAST_NAME	DEPARTMENT
1	205	Higgins	Accounting
2	206	Gietz	Accounting
3	200	Whalen	Administration
4	100	King	Executive
5	101	Kochhar	Executive
6	102	De Haan	Executive
7	109	Faviet	Finance
8	108	Greenberg	Finance
9	112	Urman	Finance
10	111	Sciarra	Finance
11	110	Chen	Finance
12	113	Popp	Finance
13	203	Mavris	Human Resources
14	107	Lorentz	IT
15	106	Pataballa	IT

...

102	140	Patel	Shipping
103	141	Rajs	Shipping
104	142	Davies	Shipping
105	143	Matos	Shipping
106	181	Fleaur	Shipping
107	178	Grant	(null)

11. Write a query to display the department names of those departments whose total salary cost is above one-eighth ($1/8$) of the total salary cost of the whole company. Use the `WITH` clause to write this query. Name the query `SUMMARY`.

	DEPARTMENT_NAME	DEPT_TOTAL
1	Sales	304500
2	Shipping	156400

Solution 6-1: Retrieving Data by Using Subqueries

Solution

1. Write a query to display the last name, department number, and salary of any employee whose department number and salary match the department number and salary of any employee who earns a commission.

```
SELECT last_name, department_id, salary
FROM   employees
WHERE  (salary, department_id) IN
        (SELECT salary, department_id
         FROM   employees
         WHERE  commission_pct IS NOT NULL);
```

2. Display the last name, department name, and salary of any employee whose salary and job_ID match the salary and job_ID of any employee located in location ID 1700.

```
SELECT e.last_name, d.department_name, e.salary
FROM   employees e JOIN departments d
ON     e.department_id = d.department_id
AND    (salary, job_id) IN
        (SELECT e.salary, e.job_id
         FROM   employees e JOIN
departments d
         ON     e.department_id =
d.department_id
         AND    d.location_id = 1700);
```

3. Create a query to display the last name, hire date, and salary for all employees who have the same salary and manager_ID as Kochhar.

Note: Do not display Kochhar in the result set.

```
SELECT last_name, hire_date, salary
FROM   employees
WHERE  (salary, manager_id) IN
        (SELECT salary, manager_id
         FROM   employees
         WHERE  last_name = 'Kochhar')
AND    last_name != 'Kochhar';
```

4. Create a query to display the employees who earn a salary that is higher than the salary of all the sales managers (JOB_ID = 'SA_MAN'). Sort the results on salary from the highest to the lowest.

```
SELECT last_name, job_id, salary
FROM   employees
WHERE  salary > ALL
        (SELECT salary
         FROM   employees
         WHERE  job_id = 'SA_MAN')
ORDER BY salary DESC;
```

5. Display details such as the employee ID, last name, and department ID of those employees who live in cities the names of which begin with T.

```
SELECT employee_id, last_name, department_id
FROM   employees
WHERE  department_id IN (SELECT department_id
                        FROM departments
                        WHERE location_id IN
                              (SELECT location_id
                               FROM locations
                               WHERE city LIKE 'T%'));
```

6. Write a query to find all employees who earn more than the average salary in their departments. Display last name, salary, department ID, and the average salary for the department. Sort by average salary and round to two decimals. Use aliases for the columns retrieved by the query as shown in the sample output.

```
SELECT e.last_name ename, e.salary salary,
       e.department_id deptno, ROUND(AVG(a.salary),2)
dept_avg
FROM   employees e, employees a
WHERE  e.department_id = a.department_id
AND    e.salary > (SELECT AVG(salary)
                  FROM   employees
                  WHERE  department_id = e.department_id )
GROUP BY e.last_name, e.salary, e.department_id
ORDER BY AVG(a.salary);
```

7. Find all employees who are not supervisors.
- a. First, do this by using the NOT EXISTS operator.

```
SELECT outer.last_name
FROM   employees outer
WHERE  NOT EXISTS (SELECT 'X'
                  FROM employees inner
                  WHERE inner.manager_id =
                        outer.employee_id);
```

- b. Can this be done by using the NOT IN operator? How, or why not?

```
SELECT outer.last_name
FROM   employees outer
WHERE  outer.employee_id
NOT IN (SELECT inner.manager_id
       FROM   employees inner);
```

This alternative solution is not a good one. The subquery picks up a NULL value, so the entire query returns no rows. The reason is that all conditions that compare a NULL value result in NULL. Whenever NULL values are likely to be part of the value set, *do not* use NOT IN as a substitute for NOT EXISTS. A much better solution would be a subquery like the following:

```
SELECT last_name
FROM employees
WHERE employee_id NOT IN (SELECT manager_id
                        FROM employees WHERE manager_id IS NOT
                        NULL);
```

8. Write a query to display the last names of the employees who earn less than the average salary in their departments.

```
SELECT last_name
FROM   employees outer
WHERE  outer.salary < (SELECT AVG(inner.salary)
                     FROM employees inner
                     WHERE inner.department_id
                           = outer.department_id);
```

9. Write a query to display the last names of employees who have one or more coworkers in their departments with later hire dates but higher salaries.

```
SELECT last_name
FROM employees outer
WHERE EXISTS (SELECT 'X'
              FROM employees inner
              WHERE inner.department_id =
                    outer.department_id
              AND inner.hire_date > outer.hire_date
              AND inner.salary > outer.salary);
```

10. Write a query to display the employee ID, last names, and department names of all employees.

Note: Use a scalar subquery to retrieve the department name in the SELECT statement.

```
SELECT employee_id, last_name,
       (SELECT department_name
        FROM departments d
        WHERE e.department_id =
              d.department_id ) department
FROM employees e
ORDER BY department;
```

11. Write a query to display the department names of those departments whose total salary cost is above one-eighth (1/8) of the total salary cost of the whole company. Use the WITH clause to write this query. Name the query SUMMARY.

```
WITH
summary AS (
  SELECT d.department_name, SUM(e.salary) AS dept_total
  FROM employees e JOIN departments d
  ON e.department_id = d.department_id
  GROUP BY d.department_name)
SELECT department_name, dept_total
FROM summary
WHERE dept_total > ( SELECT SUM(dept_total) * 1/8
                    FROM summary )
ORDER BY dept_total DESC;
```

YASSINE MAZOUZ (z.m.malak@hotmail.com) has a non-transferable license to use this Student Guide.

Practices for Lesson 7: Manipulating Data by Using Subqueries

Chapter 7

Practices for Lesson 7: Overview

Practices Overview

This practice covers the following topics:

- Using subqueries to manipulate data
- Inserting by using a subquery as a target
- Using the `WITH CHECK OPTION` keyword on DML statements
- Using correlated subqueries to update and delete rows

Practice 7-1: Manipulating Data by Using Subqueries

Overview

In this practice, you test your knowledge about using subqueries to manipulate data, using the `WITH CHECK OPTION` keyword on DML statements, and correlated subqueries to update and delete rows.

Tasks

1. Which of the following statements are true?
 - a. Subqueries are used to retrieve data by using an inline view.
 - b. Subqueries cannot be used to copy data from one table to another.
 - c. Subqueries update data in one table based on the values of another table.
 - d. Subqueries delete rows from one table based on rows in another table.
 2. Fill in the blanks:
 - a. You can use a subquery in place of the table name in the _____ clause of the `INSERT` statement.
- Options:
- 1) `FROM`
 - 2) `INTO`
 - 3) `FOR UPDATE`
 - 4) `VALUES`
3. The `WITH CHECK OPTION` keyword prohibits you from changing rows that are not in the subquery.
 - a. `TRUE`
 - b. `FALSE`
 4. The `SELECT` list of this subquery must have the same number of columns as the column list of the `VALUES` clause.
 - a. `TRUE`
 - b. `FALSE`

5. You can use a correlated subquery to delete only those rows that also exist in another table.
 - a. TRUE
 - b. FALSE
6. To understand the concepts of `WITH CHECK OPTION` and correlated subqueries, run the demo files for this practice.

Solution 7-1: Manipulating Data by Using Subqueries

1. Which of the following statements are true?
 - a. Subqueries are used to retrieve data by using an inline view.
 - b. Subqueries cannot be used to copy data from one table to another.
 - c. Subqueries update data in one table based on the values of another table.
 - d. Subqueries delete rows from one table based on rows in another table.

Answer: a, c, and d

2. Fill in the blanks:
 - a. You can use a subquery in place of the table name in the _____ clause of the INSERT statement.

Options:

- 1) FROM
- 2) INTO
- 3) FOR UPDATE
- 4) VALUES

Answer: 2

3. The WITH CHECK OPTION keyword prohibits you from changing rows that are not in the subquery.
 - a. TRUE
 - b. FALSE

Answer: a

4. The SELECT list of this subquery must have the same number of columns as the column list of the VALUES clause.
 - a. TRUE
 - b. FALSE

Answer: a

5. You can use a correlated subquery to delete only those rows that also exist in another table.
 - a. TRUE
 - b. FALSE

Answer: a

6. To understand the concepts of `WITH CHECK OPTION` and correlated subqueries, run the demo files for this practice.

Practices for Lesson 8: Controlling User Access

Chapter 8

Practices for Lesson 8: Overview

Practice Overview:

This practice covers the following topics:

- Granting other users privileges to your table
- Modifying another user's table through the privileges granted to you

Practice 8-1: Controlling User Access

Overview

You grant query privilege on your table to another user. You learn how to control access to database objects.

Tasks

1. What privilege should a user be given to log on to the Oracle server? Is this a system privilege or an object privilege?

2. What privilege should a user be given to create tables?

3. If you create a table, who can pass along privileges to other users in your table?

4. You are the DBA. You create many users who require the same system privileges. What should you use to make your job easier?

5. What command do you use to change your password?

6. User21 is the owner of the EMP table and grants the DELETE privilege to User22 by using the WITH GRANT OPTION clause. User22 then grants the DELETE privilege on EMP to User23. User21 now finds that User23 has the privilege and revokes it from User22. Which user can now delete from the EMP table?

7. You want to grant SCOTT the privilege to update data in the DEPARTMENTS table. You also want to enable SCOTT to grant this privilege to other users. What command do you use?

To complete question 8 and the subsequent ones, you need to connect to the database by using SQL Developer. If you are already not connected, do the following to connect:

1. Click the SQL Developer desktop icon.
 2. In the Connections Navigator, use the ora21 account and the corresponding password provided by your instructor to log on to the database.
 3. Open another SQL Developer session and connect as ora22.
8. Grant another user query privilege on your table. Then, verify whether that user can use the privilege.
- Note:** For this exercise, open another SQL Developer session and connect as a different user. For example, if you are currently using ora21, open another SQL Developer session and connect as ora22. Here onwards we would refer the first SQL Developer session as Team 1 and the second SQL Developer session as Team 2.
- a. Grant another user (for example, ora22) privilege to view records in your REGIONS table. Include an option for this user to further grant this privilege to other users.

- b. Have the user query your REGIONS table.

	REGION_ID	REGION_NAME
1	1	Europe
2	2	Americas
3	3	Asia
4	4	Middle East and Africa

- c. Have the user pass on the query privilege to a third user, ora23.
d. Take back the privilege from the user who performs step b.
9. Grant another user query and data manipulation privileges on your COUNTRIES table. Make sure that the user cannot pass on these privileges to other users.
10. Take back the privileges on the COUNTRIES table granted to another user.
11. Grant another user access to your DEPARTMENTS table. Have the user grant you query access to his or her DEPARTMENTS table.
12. Query all the rows in your DEPARTMENTS table.

	DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
1	10	Administration	200	1700
2	20	Marketing	201	1800
3	30	Purchasing	114	1700
4	40	Human Resources	203	2400
5	50	Shipping	121	1500
6	60	IT	103	1400
7	70	Public Relations	204	2700
8	80	Sales	145	2500
9	90	Executive	100	1700
10	100	Finance	108	1700
11	110	Accounting	205	1700
12	120	Treasury	(null)	1700
13	130	Corporate Tax	(null)	1700
14	140	Control And Credit	(null)	1700
15	150	Shareholder Services	(null)	1700
16	160	Benefits	(null)	1700
17	170	Manufacturing	(null)	1700
18	180	Construction	(null)	1700
19	190	Contracting	(null)	1700
20	200	Operations	(null)	1700

...

13. Add a new row to your DEPARTMENTS table. Team 1 should add Education as department number 500. Team 2 should add Human Resources as department number 510. Query the other team's table.
14. Create a synonym for the other team's DEPARTMENTS table.
15. Query all the rows in the other team's DEPARTMENTS table by using your synonym.

Team 1 SELECT statement results:

	DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
15	150	Shareholder Services	(null)	1700
16	160	Benefits	(null)	1700
17	170	Manufacturing	(null)	1700
18	180	Construction	(null)	1700
19	190	Contracting	(null)	1700
20	200	Operations	(null)	1700
21	210	IT Support	(null)	1700
22	220	NOC	(null)	1700
23	230	IT Helpdesk	(null)	1700
24	240	Government Sales	(null)	1700
25	250	Retail Sales	(null)	1700
26	260	Recruiting	(null)	1700
27	270	Payroll	(null)	1700
28	510	Human Resources	(null)	(null)

Team 2 SELECT statement results:

	DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
15	150	Shareholder Services	(null)	1700
16	160	Benefits	(null)	1700
17	170	Manufacturing	(null)	1700
18	180	Construction	(null)	1700
19	190	Contracting	(null)	1700
20	200	Operations	(null)	1700
21	210	IT Support	(null)	1700
22	220	NOC	(null)	1700
23	230	IT Helpdesk	(null)	1700
24	240	Government Sales	(null)	1700
25	250	Retail Sales	(null)	1700
26	260	Recruiting	(null)	1700
27	270	Payroll	(null)	1700
28	500	Education	(null)	(null)

16. Revoke the SELECT privilege from the other team.
17. Remove the row that you inserted into the DEPARTMENTS table in step 13 and save the changes.
18. Drop the synonyms team 1 and team 2.

Solution 8-1: Controlling User Access

1. What privilege should a user be given to log on to the Oracle server? Is this a system or an object privilege?

The CREATE SESSION system privilege

2. What privilege should a user be given to create tables?

The CREATE TABLE privilege

3. If you create a table, who can pass along privileges to other users in your table?

You can, or anyone you have given those privileges to, by using WITH GRANT OPTION

4. You are the DBA. You create many users who require the same system privileges. What should you use to make your job easier?

Create a role containing the system privileges and grant the role to the users.

5. What command do you use to change your password?

The ALTER USER statement

6. User21 is the owner of the EMP table and grants DELETE privileges to User22 by using the WITH GRANT OPTION clause. User22 then grants DELETE privileges on EMP to User23. User21 now finds that User23 has the privilege and revokes it from User22. Which user can now delete data from the EMP table?

Only User21

7. You want to grant SCOTT the privilege to update data in the DEPARTMENTS table. You also want to enable SCOTT to grant this privilege to other users. What command do you use?

```
GRANT UPDATE ON departments TO scott WITH GRANT OPTION;
```

8. Grant another user query privilege on your table. Then, verify whether that user can use the privilege.

Note: For this exercise, open another SQL Developer session and connect as a different user. For example, if you are currently using ora21, open another SQL Developer session and connect as ora22. Here onwards we would refer the first SQL Developer session as Team 1 and the second SQL Developer session as Team 2.

- a. Grant another user privilege to view records in your REGIONS table. Include an option for this user to further grant this privilege to other users.

Note: Replace <team2_oraxx> with ora22, <team1_oraxx> with ora21, and <team3_oraxx> with ora23.

Team 1 executes this statement:

```
GRANT select
ON regions
TO <team2_oraxx> WITH GRANT OPTION;
```

- b. Have the user query your REGIONS table.

Team 2 executes this statement:

```
SELECT * FROM <team1_oraxx>.regions;
```

- c. Have the user pass on the query privilege to a third user, ora23.

Team 2 executes this statement.

```
GRANT select
ON <team1_oraxx>.regions
TO <team3_oraxx>;
```

- d. Take back the privilege from the user who performs step b.

Team 1 executes this statement.

```
REVOKE select
ON regions
FROM <team2_oraxx>;
```

9. Grant another user query and data manipulation privileges on your COUNTRIES table. Make sure the user cannot pass on these privileges to other users.

Team 1 executes this statement.

```
GRANT select, update, insert
ON COUNTRIES
TO <team2_oraxx>;
```

10. Take back the privileges on the COUNTRIES table granted to another user.

Team 1 executes this statement.

```
REVOKE select, update, insert ON COUNTRIES FROM <team2_oraxx>;
```

11. Grant another user access to your DEPARTMENTS table. Have the user grant you query access to his or her DEPARTMENTS table.

- a. Team 2 executes the GRANT statement.

```
GRANT select
ON departments
TO <team1_oraxx>;
```

- b. Team 1 executes the GRANT statement.

```
GRANT select
ON departments
TO <team2_oraxx>;
```

Here, <team1_oraxx> is the username of Team 1 and <team2_oraxx> is the username of Team 2.

12. Query all the rows in your DEPARTMENTS table.

```
SELECT *
FROM departments;
```

13. Add a new row to your DEPARTMENTS table. Team 1 should add Education as department number 500. Team 2 should add Human Resources as department number 510. Query the other team's table.

- a. Team 1 executes this INSERT statement.

```
INSERT INTO departments(department_id, department_name)
VALUES (500, 'Education');
COMMIT;
```

- b. Team 2 executes this INSERT statement.

```
INSERT INTO departments(department_id, department_name)
VALUES (510, 'Human Resources');
COMMIT;
```

14. Create a synonym for the other team's DEPARTMENTS table.

- a. Team 1 creates a synonym named team2.

```
CREATE SYNONYM team2
FOR <team2_oraxx>.DEPARTMENTS;
```

- b. Team 2 creates a synonym named team1.

```
CREATE SYNONYM team1
FOR <team1_oraxx>. DEPARTMENTS;
```

15. Query all the rows in the other team's DEPARTMENTS table by using your synonym.

- a. Team 1 executes this SELECT statement.

```
SELECT *
FROM team2;
```

- b. Team 2 executes this SELECT statement.

```
SELECT *
FROM team1;
```

16. Revoke the `SELECT` privilege from the other team.

a. Team 1 revokes the privilege.

```
REVOKE select
  ON departments
  FROM <team2_oraxx>;
```

b. Team 2 revokes the privilege.

```
REVOKE select
  ON departments
  FROM <team1_oraxx>;
```

17. Remove the row that you inserted into the `DEPARTMENTS` table in step 13 and save the changes.

a. Team 1 executes this `DELETE` statement.

```
DELETE FROM departments
WHERE department_id = 500;
COMMIT;
```

b. Team 2 executes this `DELETE` statement.

```
DELETE FROM departments
WHERE department_id = 510;
COMMIT;
```

18. Drop the synonyms team 1 and team 2.

```
DROP SYNONYM team1;
DROP SYNONYM team2;
```

MAINEB MAZOUZ (z.m.malak@hotmail.com) has a non-transferable
license to use this Student Guide.

Practices for Lesson 9: Manipulating Data

Chapter 9

Practices for Lesson 9: Overview

Practice overview:

This practice covers the following topics:

- Performing multitable `INSERTs`
- Performing `MERGE` operations
- Performing flashback operations
- Tracking row versions

Note: Before starting this practice, execute `/home/oracle/labs/sql2/code_ex/cleanup_scripts/cleanup_09.sql` script.

Practice 9-1: Manipulating Data

Overview

In this practice, you perform multitable `INSERT` and `MERGE` operations, flashback operation, and track row versions.

Note: Execute `cleanup_09.sql` script from `/home/oracle/labs/sql2/code_ex/cleanup_scripts/` before performing the following tasks.

Tasks

1. Run the `lab_09_01.sql` script in the lab folder to create the `SAL_HISTORY` table.
2. Display the structure of the `SAL_HISTORY` table.

```
DESC sal_history
Name          Null Type
-----
EMPLOYEE_ID    NUMBER(6)
HIRE_DATE      DATE
SALARY         NUMBER(8,2)
```

3. Run the `lab_09_03.sql` script in the lab folder to create the `MGR_HISTORY` table.
4. Display the structure of the `MGR_HISTORY` table.

```
DESC mgr_history
Name          Null Type
-----
EMPLOYEE_ID    NUMBER(6)
MANAGER_ID     NUMBER(6)
SALARY         NUMBER(8,2)
```

5. Run the `lab_09_05.sql` script in the lab folder to create the `SPECIAL_SAL` table.
6. Display the structure of the `SPECIAL_SAL` table.

```
DESC special_sal
Name          Null Type
-----
EMPLOYEE_ID    NUMBER(6)
SALARY         NUMBER(8,2)
```

7.
 - a. Write a query to do the following:
 - Retrieve details such as the employee ID, hire date, salary, and manager ID of those employees whose employee ID is less than 125 from the `EMPLOYEES` table.
 - If the salary is more than \$20,000, insert details such as the employee ID and salary into the `SPECIAL_SAL` table.
 - If the salary is less than \$20,000:
 - Insert details such as the employee ID, hire date, and salary into the `SAL_HISTORY` table
 - Insert details such as the employee ID, manager ID, and salary into the `MGR_HISTORY` table

- b. Display the records from the SPECIAL_SAL table.

	EMPLOYEE_ID	SALARY
1	100	24000

- c. Display the records from the SAL_HISTORY table.

	EMPLOYEE_ID	HIRE_DATE	SALARY
1	101	21-SEP-05	17000
2	102	13-JAN-01	17000
3	103	03-JAN-06	9000
4	104	21-MAY-07	6000
5	105	25-JUN-05	4800
6	106	05-FEB-06	4800
7	107	07-FEB-07	4200
8	108	17-AUG-02	12008
9	109	16-AUG-02	9000
10	110	28-SEP-05	8200
11	111	30-SEP-05	7700
12	112	07-MAR-06	7800
13	113	07-DEC-07	6900
14	114	07-DEC-02	11000

...

- d. Display the records from the MGR_HISTORY table.

	EMPLOYEE_ID	MANAGER_ID	SALARY
1	101	100	17000
2	102	100	17000
3	103	102	9000
4	104	103	6000
5	105	103	4800
6	106	103	4800
7	107	103	4200
8	108	101	12008
9	109	108	9000
10	110	108	8200
11	111	108	7700
12	112	108	7800
13	113	108	6900

...

8.

- a. Run the lab_09_08_a.sql script in the lab folder to create the SALES_WEEK_DATA table.
- b. Run the lab_09_08_b.sql script in the lab folder to insert records into the SALES_WEEK_DATA table.
- c. Display the structure of the SALES_WEEK_DATA table.

```
DESC sales_week_data
```

Name	Null	Type
ID		NUMBER(6)
WEEK_ID		NUMBER(2)
QTY_MON		NUMBER(8,2)
QTY_TUE		NUMBER(8,2)
QTY_WED		NUMBER(8,2)
QTY_THUR		NUMBER(8,2)
QTY_FRI		NUMBER(8,2)

- d. Display the records from the SALES_WEEK_DATA table.

	ID	WEEK_ID	QTY_MON	QTY_TUE	QTY_WED	QTY_THUR	QTY_FRI
1	200	6	2050	2200	1700	1200	3000

- e. Run the lab_09_08_e.sql script in the lab folder to create the EMP_SALES_INFO table.
- f. Display the structure of the EMP_SALES_INFO table.

```
DESC emp_sales_info
```

Name	Null	Type
ID		NUMBER(6)
WEEK		NUMBER(2)
QTY_SALES		NUMBER(8,2)

- g. Write a query to do the following:
 - Retrieve details such as ID, week ID, sales quantity on Monday, sales quantity on Tuesday, sales quantity on Wednesday, sales quantity on Thursday, and sales quantity on Friday from the SALES_WEEK_DATA table.
 - Build a transformation such that each record retrieved from the SALES_WEEK_DATA table is converted into multiple records for the EMP_SALES_INFO table.

Hint: Use a pivoting INSERT statement.

- h. Display the records from the EMP_SALES_INFO table.

R	ID	WEEK	QTY_SALES
1	200	6	2050
2	200	6	2200
3	200	6	1700
4	200	6	1200
5	200	6	3000

9. You have the data of past employees stored in a flat file called `emp.dat`. You want to store the names and email IDs of all employees, past and present, in a table. To do this, first create an external table called `EMP_DATA` using the `emp.dat` source file in the `emp_dir` directory. Use the `lab_09_09.sql` script to do this.
10. Run the `lab_09_10.sql` script to create the `EMP_HIST` table.
- Increase the size of the email column to 45.
 - Merge the data in the `EMP_DATA` table created in the last lab into the data in the `EMP_HIST` table. Assume that the data in the external `EMP_DATA` table is the most up-to-date. If a row in the `EMP_DATA` table matches the `EMP_HIST` table, update the email column of the `EMP_HIST` table to match the `EMP_DATA` table row. If a row in the `EMP_DATA` table does not match, insert it into the `EMP_HIST` table. Rows are considered matching when the employee's first and last names are identical.
 - Retrieve the rows from `EMP_HIST` after the merge.

R	FIRST_NAME	LAST_NAME	EMAIL
1	Ellen	Abel	EABEL
2	Sundar	Ande	SANDE
3	Mozhe	Atkinson	MATKINSO
4	David	Austin	DAUSTIN
5	Hermann	Baer	HBAER
6	Shelli	Baida	SBIDA
7	Amit	Banda	ABANDA
8	Elizabeth	Bates	EBATES
9	Sarah	Bell	SBELL
10	David	Bernstein	DBERNSTE
11	Laura	Bissot	LBISSOT
12	Harrison	Bloom	HBLOOM

...

11. Create the EMP2 table based on the following table instance chart. Enter the syntax in the SQL Worksheet. Then execute the statement to create the table. Confirm that the table is created.

Column Name	ID	LAST_NAME	FIRST_NAME	DEPT_ID
Key Type				
Nulls/Unique				
FK Table				
FK Column				
Data type	NUMBER	VARCHAR2	VARCHAR2	NUMBER
Length	7	25	25	7

12. Drop the EMP2 table.
13. Query the recycle bin to see whether the table is present.
14. Restore the EMP2 table to a state before the DROP statement.
15. Create the EMP3 table using the lab_09_11.sql script. In the EMP3 table, change the department for Kochhar to 60 and commit your change. Next, change the department for Kochhar to 50 and commit your change. Track the changes to Kochhar using the Row Versions feature.

```

UPDATE emp3 SET department_id = 60
WHERE last_name = 'Kochhar';
COMMIT;
UPDATE emp3 SET department_id = 50
WHERE last_name = 'Kochhar';
COMMIT;

SELECT VERSIONS_STARTTIME "START_DATE",
       VERSIONS_ENDTIME "END_DATE",  DEPARTMENT_ID
FROM EMP3
      VERSIONS BETWEEN SCN MINVALUE AND MAXVALUE
WHERE LAST_NAME = 'Kochhar';

```

	START_DATE	END_DATE	DEPARTMENT_ID
1	28-APR-14 10.11.40.000000000 PM (null)		50
2	28-APR-14 10.11.37.000000000 PM	28-APR-14 10.11.40.000000000 PM	60
3	(null)	28-APR-14 10.11.37.000000000 PM	90

16. Drop the EMP2 and EMP3 tables so that they cannot be restored. Check in the recycle bin.

Solution 9-1: Manipulating Data

Solution

1. Run the `lab_09_01.sql` script in the lab folder to create the `SAL_HISTORY` table.
2. Display the structure of the `SAL_HISTORY` table.

```
DESC sal_history
```

3. Run the `lab_09_03.sql` script in the lab folder to create the `MGR_HISTORY` table.
4. Display the structure of the `MGR_HISTORY` table.

```
DESC mgr_history
```

5. Run the `lab_09_05.sql` script in the lab folder to create the `SPECIAL_SAL` table.
6. Display the structure of the `SPECIAL_SAL` table.

```
DESC special_sal
```

7.
 - a. Write a query to do the following:
 - Retrieve details such as the employee ID, hire date, salary, and manager ID of those employees whose employee ID is less than 125 from the `EMPLOYEES` table.
 - If the salary is more than \$20,000, insert details such as the employee ID and salary into the `SPECIAL_SAL` table.
 - If the salary is less than \$20,000:
 - Insert details such as the employee ID, hire date, and salary into the `SAL_HISTORY` table
 - Insert details such as the employee ID, manager ID, and salary into the `MGR_HISTORY` table

```
INSERT ALL
WHEN SAL > 20000 THEN
  INTO special_sal VALUES (EMPID, SAL)
ELSE
  INTO sal_history VALUES (EMPID, HIREDATE, SAL)
  INTO mgr_history VALUES (EMPID, MGR, SAL)
SELECT employee_id EMPID, hire_date HIREDATE,
       salary SAL, manager_id MGR
FROM employees
WHERE employee_id < 125;
```


- b. Display the records from the SPECIAL_SAL table.

```
SELECT * FROM special_sal;
```

- c. Display the records from the SAL_HISTORY table.

```
SELECT * FROM sal_history;
```

- d. Display the records from the MGR_HISTORY table.

```
SELECT * FROM mgr_history;
```

8.

- a. Run the lab_09_08_a.sql script in the lab folder to create the SALES_WEEK_DATA table.
- b. Run the lab_09_08_b.sql script in the lab folder to insert records into the SALES_WEEK_DATA table.
- c. Display the structure of the SALES_WEEK_DATA table.

```
DESC sales_week_data
```

- d. Display the records from the SALES_WEEK_DATA table.

```
SELECT * FROM SALES_WEEK_DATA;
```

- e. Run the lab_09_08_e.sql script in the lab folder to create the EMP_SALES_INFO table.
- f. Display the structure of the EMP_SALES_INFO table.

```
DESC emp_sales_info
```

- g. Write a query to do the following:

- Retrieve details such as the employee ID, week ID, sales quantity on Monday, sales quantity on Tuesday, sales quantity on Wednesday, sales quantity on Thursday, and sales quantity on Friday from the SALES_WEEK_DATA table.
- Build a transformation such that each record retrieved from the SALES_WEEK_DATA table is converted into multiple records for the EMP_SALES_INFO table.

Hint: Use a pivoting INSERT statement.

```
INSERT ALL
    INTO emp_sales_info VALUES (id, week_id, QTY_MON)
    INTO emp_sales_info VALUES (id, week_id, QTY_TUE)
    INTO emp_sales_info VALUES (id, week_id, QTY_WED)
    INTO emp_sales_info VALUES (id, week_id, QTY_THUR)
    INTO emp_sales_info VALUES (id, week_id, QTY_FRI)
```

```
SELECT ID, week_id, QTY_MON, QTY_TUE, QTY_WED,  
       QTY_THUR,QTY_FRI FROM sales_week_data;
```

- h. Display the records from the SALES_INFO table.

```
SELECT * FROM emp_sales_info;
```

9. You have the data of past employees stored in a flat file called `emp.data`. You want to store the names and email IDs of all employees past and present in a table. To do this, first create an external table called `EMP_DATA` using the `emp.dat` source file in the `emp_dir` directory. You can use the script in `lab_09_09.sql` to do this.

```
CREATE TABLE emp_data  
(first_name VARCHAR2(20)  
,last_name VARCHAR2(20)  
, email VARCHAR2(30)  
)  
ORGANIZATION EXTERNAL  
(  
  TYPE oracle_loader  
  DEFAULT DIRECTORY emp_dir  
  ACCESS PARAMETERS  
  (  
    RECORDS DELIMITED BY NEWLINE CHARACTERSET US7ASCII  
    NOBADFILE  
    NOLOGFILE  
    FIELDS  
    ( first_name POSITION ( 1:20) CHAR  
    , last_name POSITION (22:41) CHAR  
    , email POSITION (43:72) CHAR )  
  )  
  LOCATION ('emp.dat') ) ;
```

10. Run the `lab_09_10.sql` script to create the `EMP_HIST` table.

- a. Increase the size of the email column to 45.

```
ALTER TABLE emp_hist MODIFY email varchar(45);
```

- b. Merge the data in the `EMP_DATA` table created in the last lab into the data in the `EMP_HIST` table. Assume that the data in the external `EMP_DATA` table is the most up-to-date. If a row in the `EMP_DATA` table matches the `EMP_HIST` table, update the email column of the `EMP_HIST` table to match the `EMP_DATA` table row. If a row in the

EMP_DATA table does not match, insert it into the EMP_HIST table. Rows are considered matching when the employee's first and last names are identical.

```

MERGE INTO EMP_HIST f USING EMP_DATA h
  ON (f.first_name = h.first_name
     AND f.last_name = h.last_name)
WHEN MATCHED THEN
  UPDATE SET f.email = h.email
WHEN NOT MATCHED THEN
  INSERT (f.first_name
        , f.last_name
        , f.email)
  VALUES (h.first_name
        , h.last_name
        , h.email);

```

- c. Retrieve the rows from EMP_HIST after the merge.

```
SELECT * FROM emp_hist;
```

11. Create the EMP2 table based on the following table instance chart. Enter the syntax in the SQL Worksheet. Then execute the statement to create the table. Confirm that the table is created.

Column Name	ID	LAST_NAME	FIRST_NAME	DEPT_ID
Key Type				
Nulls/Unique				
FK Table				
FK Column				
Data type	NUMBER	VARCHAR2	VARCHAR2	NUMBER
Length	7	25	25	7

```
CREATE TABLE emp2
(id          NUMBER(7),
last_name    VARCHAR2(25),
first_name   VARCHAR2(25),
dept_id      NUMBER(7));

DESCRIBE emp2
```

12. Drop the EMP2 table.

```
DROP TABLE emp2;
```

13. Query the recycle bin to see whether the table is present.

```
SELECT original_name, operation, droptime
FROM recyclebin;
```

14. Restore the EMP2 table to a state before the DROP statement.

```
FLASHBACK TABLE emp2 TO BEFORE DROP;
DESC emp2;
```

15. Create the EMP3 table using the lab_09_11.sql script. In the EMP3 table, change the department for Kochhar to 60 and commit your change. Next, change the department for Kochhar to 50 and commit your change. Track the changes to Kochhar using the Row Versions feature.

```
UPDATE emp3 SET department_id = 60
WHERE last_name = 'Kochhar';
COMMIT;

UPDATE emp3 SET department_id = 50
WHERE last_name = 'Kochhar';
COMMIT;

SELECT VERSIONS_STARTTIME "START_DATE",
       VERSIONS_ENDTIME "END_DATE",  DEPARTMENT_ID
FROM EMP3
      VERSIONS BETWEEN SCN MINVALUE AND MAXVALUE
WHERE LAST_NAME = 'Kochhar';
```

16. Drop the EMP2 and EMP3 tables so that they cannot be restored. Check in the recycle bin.

```
DROP TABLE emp2 PURGE;  
DROP TABLE emp3 PURGE;  
  
SELECT original_name, operation, droptime  
FROM recyclebin;
```

YASSINE MAZOUZ (z.m.malak@hotmail.com) has a non-transferable license to use this Student Guide.

Practices for Lesson 10: Managing Data in Different Time Zones

Chapter 10

Practices for Lesson 10: Overview

Practice Overview:

This practice covers using the datetime functions.

Note: Before starting this practice, execute
`/home/oracle/labs/sql2/code_ex/cleanup_scripts/cleanup_10.sql` script.

Practice 10-1: Managing Data in Different Time Zones

Overview

In this practice, you display time zone offsets, `CURRENT_DATE`, `CURRENT_TIMESTAMP`, and `LOCALTIMESTAMP`. You also set time zones and use the `EXTRACT` function.

Note: Execute `cleanup_10.sql` script from `/home/oracle/labs/sql2/code_ex/cleanup_scripts/cleanup_10.sql` before performing the following tasks.

Tasks

1. Alter the session to set `NLS_DATE_FORMAT` to `DD-MON-YYYY HH24:MI:SS`.
2.
 - a. Write queries to display the time zone offsets (`TZ_OFFSET`) for the following time zones.

- US/Pacific-New

	TZ_OFFSET('US/PACIFIC-NEW')
1	-07:00

- Singapore

	TZ_OFFSET('SINGAPORE')
1	+08:00

- Egypt

	TZ_OFFSET('EGYPT')
1	+02:00

- b. Alter the session to set the `TIME_ZONE` parameter value to the time zone offset of US/Pacific-New.
- c. Display `CURRENT_DATE`, `CURRENT_TIMESTAMP`, and `LOCALTIMESTAMP` for this session.

	CURRENT_DATE	CURRENT_TIMESTAMP	LOCALTIMESTAMP
1	16-SEP-2012 21:03:47	16-SEP-12 09.03.47.344991000 PM -07:00	16-SEP-12 09.03.47.344991000 PM

- d. Alter the session to set the `TIME_ZONE` parameter value to the time zone offset of Singapore.
- e. Display `CURRENT_DATE`, `CURRENT_TIMESTAMP`, and `LOCALTIMESTAMP` for this session.

Note: The output might be different based on the date when the command is executed.

	CURRENT_DATE	CURRENT_TIMESTAMP	LOCALTIMESTAMP
1	17-SEP-2012 12:05:56	17-SEP-12 12.05.56.424157000 PM +08:00	17-SEP-12 12.05.56.424157000 PM

Note: Observe in the preceding practice that `CURRENT_DATE`, `CURRENT_TIMESTAMP`, and `LOCALTIMESTAMP` are sensitive to the session time zone.

- Write a query to display DBTIMEZONE and SESSIONTIMEZONE.

	DBTIMEZONE	SESSIONTIMEZONE
1	+00:00	+08:00

- Write a query to extract the YEAR from the HIRE_DATE column of the EMPLOYEES table for those employees who work in department 80.

	LAST_NAME	EXTRACT(YEARFROMHIRE_DATE)
1	Russell	2004
2	Partners	2005
3	Errazuriz	2005
4	Cambrault	2007
5	Zlotkey	2008
6	Tucker	2005
7	Bernstein	2005
8	Hall	2005
9	Olsen	2006
10	Cambrault	2006
11	Tuvault	2007

...

- Alter the session to set NLS_DATE_FORMAT to DD-MON-YYYY.
- Examine and run the lab_10_06.sql script to create the SAMPLE_DATES table and populate it.

Note: The screenshot dates will change according to the sysdate.

- Select from the table and view the data.

	DATE_COL
1	17-SEP-2012

- Modify the data type of the DATE_COL column and change it to TIMESTAMP. Select from the table to view the data.

	DATE_COL
1	17-SEP-12 04.09.12.000000000 AM

- Try to modify the data type of the DATE_COL column and change it to TIMESTAMP WITH TIME ZONE. What happens?

```
Error report:
SQL Error: ORA-01439: column to be modified must be empty to change datatype
01439. 00000 - "column to be modified must be empty to change datatype"
*Cause:
*Action:
```

7. Create a query to retrieve last names from the EMPLOYEES table and calculate the review status. If the year hired was 2008, display Needs Review for the review status; otherwise, display not this year! Name the review status column Review. Sort the results by the HIRE_DATE column.

Hint: Use a CASE expression with the EXTRACT function to calculate the review status.

	LAST_NAME	Review
1	De Haan	not this year!
2	Gietz	not this year!
3	Baer	not this year!
4	Mavris	not this year!
5	Higgins	not this year!
6	Faviet	not this year!
7	Greenberg	not this year!
8	Raphaely	not this year!
9	Kaufling	not this year!

...

8. Create a query to print the last names and the number of years of service for each employee. If the employee has been employed for five or more years, print 5 years of service. If the employee has been employed for 10 or more years, print 10 years of service. If the employee has been employed for 15 or more years, print 15 years of service. If none of these conditions matches, print maybe next year! Sort the results by the HIRE_DATE column. Use the EMPLOYEES table.

Hint: Use CASE expressions and TO_YMINTERVAL.

	LAST_NAME	HIRE_DATE	SYSDATE	Awards
1	King	17-JUN-03	17-SEP-2012	5 years of service
2	Kochhar	21-SEP-05	17-SEP-2012	5 years of service
3	De Haan	13-JAN-01	17-SEP-2012	10 years of service
4	Hunold	03-JAN-06	17-SEP-2012	5 years of service
5	Ernst	21-MAY-07	17-SEP-2012	5 years of service
6	Austin	25-JUN-05	17-SEP-2012	5 years of service
7	Pataballa	05-FEB-06	17-SEP-2012	5 years of service
8	Lorentz	07-FEB-07	17-SEP-2012	5 years of service
9	Greenberg	17-AUG-02	17-SEP-2012	10 years of service
10	Faviet	16-AUG-02	17-SEP-2012	10 years of service
11	Chen	28-SEP-05	17-SEP-2012	5 years of service
12	Sciarra	30-SEP-05	17-SEP-2012	5 years of service
13	Urman	07-MAR-06	17-SEP-2012	5 years of service

...

Solution 10-1: Managing Data in Different Time Zones

Solution

1. Alter the session to set NLS_DATE_FORMAT to DD-MON-YYYY HH24:MI:SS.

```
ALTER SESSION SET NLS_DATE_FORMAT =  
'DD-MON-YYYY HH24:MI:SS';
```

2.

- a. Write queries to display the time zone offsets (TZ_OFFSET) for the following time zones: *US/Pacific-New*, *Singapore*, and *Egypt*.

```
US/Pacific-New  
SELECT TZ_OFFSET ('US/Pacific-New') from dual;  
Singapore  
SELECT TZ_OFFSET ('Singapore') from dual;  
Egypt  
SELECT TZ_OFFSET ('Egypt') from dual;
```

- b. Alter the session to set the TIME_ZONE parameter value to the time zone offset of US/Pacific-New.

```
ALTER SESSION SET TIME_ZONE = '-7:00';
```

- c. Display CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP for this session.

Note: The output may be different based on the date when the command is executed.

```
SELECT CURRENT_DATE, CURRENT_TIMESTAMP,  
LOCALTIMESTAMP FROM DUAL;
```

- d. Alter the session to set the TIME_ZONE parameter value to the time zone offset of Singapore.

```
ALTER SESSION SET TIME_ZONE = '+8:00';
```

- e. Display CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP for this session.

Note: The output might be different, based on the date when the command is executed.

```
SELECT CURRENT_DATE, CURRENT_TIMESTAMP,  
LOCALTIMESTAMP FROM DUAL;
```

Note: Observe in the preceding practice that CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP are all sensitive to the session time zone.

3. Write a query to display DBTIMEZONE and SESSIONTIMEZONE.

```
SELECT DBTIMEZONE, SESSIONTIMEZONE
FROM DUAL;
```

4. Write a query to extract YEAR from the HIRE_DATE column of the EMPLOYEES table for those employees who work in department 80.

```
SELECT last_name, EXTRACT (YEAR FROM HIRE_DATE)
FROM employees
WHERE department_id = 80;
```

5. Alter the session to set NLS_DATE_FORMAT to DD-MON-YYYY.

```
ALTER SESSION SET NLS_DATE_FORMAT = 'DD-MON-YYYY';
```

6. Examine and run the lab_10_06.sql script to create the SAMPLE_DATES table and populate it.

- a. Select from the table and view the data.

```
SELECT * FROM sample_dates;
```

- b. Modify the data type of the DATE_COL column and change it to TIMESTAMP. Select from the table to view the data.

```
ALTER TABLE sample_dates MODIFY date_col TIMESTAMP;
SELECT * FROM sample_dates;
```

- c. Try to modify the data type of the DATE_COL column and change it to TIMESTAMP WITH TIME ZONE. What happens?

```
ALTER TABLE sample_dates MODIFY date_col
TIMESTAMP WITH TIME ZONE;
```

You are unable to change the data type of the DATE_COL column because the Oracle server does not permit you to convert from TIMESTAMP to TIMESTAMP WITH TIMEZONE by using the ALTER statement.

7. Create a query to retrieve last names from the EMPLOYEES table and calculate the review status. If the year hired was 2008, display Needs Review for the review status; otherwise, display not this year! Name the review status column Review. Sort the results by the HIRE_DATE column.

Hint: Use a CASE expression with the EXTRACT function to calculate the review status.

```
SELECT e.last_name
       ,      (CASE extract(year from e.hire_date)
                WHEN 2008 THEN 'Needs Review'
                ELSE 'not this year!'
              END )          AS "Review "
FROM   employees e
ORDER BY e.hire_date;
```

8. Create a query to print the last names and the number of years of service for each employee. If the employee has been employed five or more years, print 5 years of service. If the employee has been employed 10 or more years, print 10 years of service. If the employee has been employed 15 or more years, print 15 years of service. If none of these conditions matches, print maybe next year! Sort the results by the HIRE_DATE column. Use the EMPLOYEES table.

Hint: Use CASE expressions and TO_YMINTERVAL.

```
SELECT e.last_name, hire_date, sysdate,
       (CASE
        WHEN (sysdate -TO_YMINTERVAL('15-0'))>=
             hire_date THEN '15 years of service'
        WHEN (sysdate -TO_YMINTERVAL('10-0'))>= hire_date
             THEN '10 years of service'
        WHEN (sysdate - TO_YMINTERVAL('5-0'))>= hire_date
             THEN '5 years of service'
        ELSE 'maybe next year!'
       END) AS "Awards"
FROM   employees e;
```

Additional Practices and Solutions

Chapter 11

Additional Practices and Solutions

Practices Overview:

In these practices, you will be working on extra practices that are based on the following topics:

- Data manipulation language (DML) statements
- Data definition language (DDL) statements
- Datetime functions
- Advanced subqueries

Additional Practices

Overview

The following exercises can be used for extra practice after you have discussed data manipulation language (DML) and data definition language (DDL) statements in the lessons titled “Managing Schema Objects” and “Manipulating Data.”

Note: Run the `lab_ap_cre_special_sal.sql`, `lab_ap_cre_sal_history.sql`, and `lab_ap_cre_mgr_history.sql` scripts in the labs folder to create the `SPECIAL_SAL`, `SAL_HISTORY`, and `MGR_HISTORY` tables.

Tasks

1. The Human Resources department wants to get a list of underpaid employees, salary history of employees, and salary history of managers based on an industry salary survey. So they have asked you to do the following:

Write a statement to do the following:

- Retrieve details such as the employee ID, hire date, salary, and manager ID of those employees whose employee ID is more than or equal to 200 from the `EMPLOYEES` table.
- If the salary is less than \$5,000, insert details such as the employee ID and salary into the `SPECIAL_SAL` table.
- Insert details such as the employee ID, hire date, and salary into the `SAL_HISTORY` table.
- Insert details such as the employee ID, manager ID, and salary into the `MGR_HISTORY` table.

2. Query the `SPECIAL_SAL`, `SAL_HISTORY`, and `MGR_HISTORY` tables to view the inserted records.

SPECIAL_SAL

	EMPLOYEE_ID	SALARY
1	200	4400

SAL_HISTORY

	EMPLOYEE_ID	HIRE_DATE	SALARY
1	201	17-FEB-04	13000
2	202	17-AUG-05	6000
3	203	07-JUN-02	6500
4	204	07-JUN-02	10000
5	205	07-JUN-02	12008
6	206	07-JUN-02	8300

MGR_HISTORY

	EMPLOYEE_ID	MANAGER_ID	SALARY
1	201	100	13000
2	202	201	6000
3	203	101	6500
4	204	101	10000
5	205	101	12008
6	206	205	8300

3. Nita, the DBA, needs you to create a table, which has a primary key constraint, but she wants the index to have a different name than the constraint. Create the LOCATIONS_NAMED_INDEX table based on the following table instance chart. Name the index for the PRIMARY KEY column as LOCATIONS_PK_IDX.

Column Name	Deptno	Dname
Primary Key	Yes	
Data Type	Number	VARCHAR2
Length	4	30

4. Query the USER_INDEXES table to display the INDEX_NAME for the LOCATIONS_NAMED_INDEX table.

	INDEX_NAME	TABLE_NAME
1	LOCATIONS_PK_IDX	LOCATIONS_NAMED_INDEX

The following exercises can be used for extra practice after you have discussed datetime functions.

You work for a global company and the new vice president of operations wants to know the different time zones of all the company branches. The new vice president has requested the following information:

5. Alter the session to set the NLS_DATE_FORMAT to DD-MON-YYYY HH24:MI:SS.
6.
 - a. Write queries to display the time zone offsets (TZ_OFFSET) for the following time zones:
 - Australia/Sydney

	TZ_OFFSET('AUSTRALIA/SYDNEY')
1	+10:00

- Chile/Easter Island

	TZ_OFFSET('CHILE/EASTERISLAND')
1	-06:00

- b. Alter the session to set the TIME_ZONE parameter value to the time zone offset of Australia/Sydney.
- c. Display SYSDATE, CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP for this session.

Note: The output might be different based on the date when the command is executed.

	SYSDATE		CURRENT_DATE		CURRENT_TIMESTAMP		LOCALTIMESTAMP
1	17-SEP-2012 22:59:31	18-SEP-2012 08:59:31	18-SEP-12 08.59.31.390170000	AM	+10:00	18-SEP-12 08.59.31.	

- d. Alter the session to set the TIME_ZONE parameter value to the time zone offset of Chile/Easter Island.

Note: The results of the preceding question are based on a different date, and in some cases, they will not match the actual results that the students get. In addition, the time zone offset of the various countries may differ, based on daylight saving time.
- e. Display SYSDATE, CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP for this session.

Note: The output may be different based on the date when the command is executed.

	SYSDATE		CURRENT_DATE		CURRENT_TIMESTAMP		LOCALTIMESTAMP
1	17-SEP-2012 23:00:01	17-SEP-2012 17:00:01	17-SEP-12 05.00.01.121124000	PM	-06:00	17-SEP-12 05.00.01.	

- f. Alter the session to set NLS_DATE_FORMAT to DD-MON-YYYY.

Note

- Observe in the preceding question that CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP are all sensitive to the session time zone. Observe that SYSDATE is not sensitive to the session time zone.
- The results of the preceding question are based on a different date, and in some cases, they will not match the actual results that the students get. In addition, the time zone offset of the various countries may differ, based on daylight saving time.

7. The Human Resources department wants a list of employees who are up for review in January, so they have requested you to do the following:

Write a query to display the last names, month of the date of hire, and hire date of those employees who have been hired in the month of January, irrespective of the year of hire.

	LAST_NAME	EXTRACT(MONTHFROMHIRE_DATE)	HIRE_DATE
1	De Haan		13-JAN-2001
2	Hunold		03-JAN-2006
3	Landry		14-JAN-2007
4	Davies		29-JAN-2005
5	Partners		05-JAN-2005
6	Zlotkey		29-JAN-2008
7	Tucker		30-JAN-2005
8	King		30-JAN-2004
9	Marvins		24-JAN-2008
10	Fox		24-JAN-2006
11	Johnson		04-JAN-2008
12	Taylor		24-JAN-2006
13	Sarchand		27-JAN-2004
14	Grant		13-JAN-2008

The following exercises can be used for extra practice after you have discussed advanced subqueries.

8. The CEO needs a report on the top three earners in the company for profit sharing. You are responsible to provide the CEO with a list. Write a query to display the top three earners in the `EMPLOYEES` table. Display their last names and salaries.

	LAST_NAME	SALARY
1	King	24000
2	Kochhar	17000
3	De Haan	17000

9. The benefits for the state of California have been changed based on a local ordinance. So the benefits representative has asked you to compile a list of the people who are affected. Write a query to display the employee ID and last names of the employees who work in the state of California.

Hint: Use scalar subqueries.

	EMPLOYEE_ID	LAST_NAME
1	120	Weiss
2	121	Fripp
3	122	Kaufling
4	123	Vollman
5	124	Mourgos
6	125	Nayer
7	126	Mikkilineni
8	127	Landry
9	128	Markle
10	129	Bissot
11	130	Atkinson
12	131	Marlow
13	132	Olson
14	133	Mallin
15	134	Rogers
16	135	Gee
17	136	Philtanker
18	137	Ladwig

...

10. Nita, the DBA, wants to remove old information from the database. One of the things she thinks is unnecessary is the old employment records. She has asked you to do the following:

Write a query to delete the oldest `JOB_HISTORY` row of an employee by looking up the `JOB_HISTORY` table for the `MIN (START_DATE)` for the employee. Delete the records of *only* those employees who have changed at least two jobs.

Hint: Use a correlated `DELETE` command.

11. The vice president of Human Resources needs the complete employment records for the annual employee recognition banquet speech. The vice president makes a quick phone call to stop you from following the DBA's orders.

Roll back the transaction.

12. The sluggish economy is forcing management to take cost reduction actions. The CEO wants to review the highest paid jobs in the company. You are responsible to provide the CEO with a list based on the following specifications:

Write a query to display the job IDs of those jobs whose maximum salary is above half the maximum salary in the entire company. Use the `WITH` clause to write this query. Name the query `MAX_SAL_CALC`.

R2	JOB_TITLE	R2	JOB_TOTAL
1	President		24000
2	Administration Vice President		17000
3	Sales Manager		14000
4	Marketing Manager		13000
5	Finance Manager		12008
6	Accounting Manager		12008

Additional Practices Solutions

Solution

The following exercises can be used for extra practice after you have discussed data manipulation language (DML) and data definition language (DDL) statements in the lessons titled “Managing Schema Objects” and “Manipulating Data.”

Note: Run the `lab_ap_cre_special_sal.sql`, `lab_ap_cre_sal_history.sql`, and `lab_ap_cre_mgr_history.sql` scripts in the labs folder to create the `SPECIAL_SAL`, `SAL_HISTORY`, and `MGR_HISTORY` tables

1. The Human Resources department wants to get a list of underpaid employees, salary history of employees, and salary history of managers based on an industry salary survey. So they have asked you to do the following:

Write a statement to do the following:

- Retrieve details such as the employee ID, hire date, salary, and manager ID of those employees whose employee ID is more than or equal to 200 from the `EMPLOYEES` table.
- If the salary is less than \$5,000, insert details such as the employee ID and salary into the `SPECIAL_SAL` table.
- Insert details such as the employee ID, hire date, and salary into the `SAL_HISTORY` table.
- Insert details such as the employee ID, manager ID, and salary into the `MGR_HISTORY` table.

```
INSERT ALL
WHEN SAL < 5000 THEN
INTO special_sal VALUES (EMPID, SAL)
ELSE
INTO sal_history VALUES (EMPID, HIREDATE, SAL)
INTO mgr_history VALUES (EMPID, MGR, SAL)
SELECT employee_id EMPID, hire_date HIREDATE,
       salary SAL, manager_id MGR
FROM employees
WHERE employee_id >= 200;
```

2. Query the `SPECIAL_SAL`, `SAL_HISTORY`, and the `MGR_HISTORY` tables to view the inserted records.

```
SELECT * FROM special_sal;
SELECT * FROM sal_history;
SELECT * FROM mgr_history;
```

3. Nita, the DBA, needs you to create a table, which has a primary key constraint, but she wants the index to have a different name than the constraint. Create the `LOCATIONS_NAMED_INDEX` table based on the following table instance chart. Name the index for the `PRIMARY KEY` column as `LOCATIONS_PK_IDX`.

Column Name	Deptno	Dname
Primary Key	Yes	
Data Type	Number	VARCHAR2
Length	4	30

```
CREATE TABLE LOCATIONS_NAMED_INDEX
(location_id NUMBER(4) PRIMARY KEY USING INDEX
(CREATE INDEX locations_pk_idx ON
LOCATIONS_NAMED_INDEX(location_id),
location_name VARCHAR2(20));
```

4. Query the `USER_INDEXES` table to display the `INDEX_NAME` for the `LOCATIONS_NAMED_INDEX` table.

```
SELECT INDEX_NAME, TABLE_NAME
FROM USER_INDEXES
WHERE TABLE_NAME = 'LOCATIONS_NAMED_INDEX';
```


The following exercises can be used for extra practice after you have discussed datetime functions.

You work for a global company and the new vice president of operations wants to know the different time zones of all the company branches. The new vice president has requested the following information:

5. Alter the session to set NLS_DATE_FORMAT to DD-MON-YYYY HH24:MI:SS.

```
ALTER SESSION
SET NLS_DATE_FORMAT = 'DD-MON-YYYY HH24:MI:SS';
```

6.

- a. Write queries to display the time zone offsets (TZ_OFFSET) for the following time zones:

- Australia/Sydney

```
SELECT TZ_OFFSET ('Australia/Sydney') from dual;
```

- Chile/Easter Island

```
SELECT TZ_OFFSET ('Chile/EasterIsland') from dual;
```

- b. Alter the session to set the TIME_ZONE parameter value to the time zone offset of Australia/Sydney.

```
ALTER SESSION SET TIME_ZONE = '+10:00';
```

- c. Display SYSDATE, CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP for this session.

Note: The output may be different based on the date when the command is executed.

```
SELECT SYSDATE, CURRENT_DATE, CURRENT_TIMESTAMP,
LOCALTIMESTAMP FROM DUAL;
```

- d. Alter the session to set the TIME_ZONE parameter value to the time zone offset of Chile/Easter Island.

Note: The results of the preceding question are based on a different date, and in some cases, they will not match the actual results that the students get. In addition, the time zone offset of the various countries may differ, based on daylight saving time.

```
ALTER SESSION SET TIME_ZONE = '-06:00';
```

- e. Display SYSDATE, CURRENT_DATE, CURRENT_TIMESTAMP, and LOCALTIMESTAMP for this session.

Note: The output may be different based on the date when the command is executed.

```
SELECT SYSDATE, CURRENT_DATE, CURRENT_TIMESTAMP,
LOCALTIMESTAMP FROM DUAL;
```

- f. Alter the session to set NLS_DATE_FORMAT to DD-MON-YYYY.

```
ALTER SESSION SET NLS_DATE_FORMAT = 'DD-MON-YYYY';
```

Note

- Observe in the preceding question that `CURRENT_DATE`, `CURRENT_TIMESTAMP`, and `LOCALTIMESTAMP` are all sensitive to the session time zone. Observe that `SYSDATE` is not sensitive to the session time zone.
- The results of the preceding question are based on a different date, and in some cases, they will not match the actual results that the students get. In addition, the time zone offset of the various countries may differ, based on daylight saving time.

7. The Human Resources department wants a list of employees who are up for review in January, so they have requested you to do the following:

Write a query to display the last names, month of the date of hire, and hire date of those employees who have been hired in the month of January, irrespective of the year of hire.

```
SELECT last_name, EXTRACT (MONTH FROM HIRE_DATE), HIRE_DATE
FROM employees
WHERE EXTRACT (MONTH FROM HIRE_DATE) = 1;
```

The following exercises can be used for extra practice after you have discussed advanced subqueries.

8. The CEO needs a report on the top three earners in the company for profit sharing. You are responsible to provide the CEO with a list. Write a query to display the top three earners in the `EMPLOYEES` table. Display their last names and salaries.

```
SELECT last_name, salary
FROM employees e
WHERE 3 > (SELECT COUNT (*)
FROM employees
WHERE e.salary < salary);
```

9. The benefits for the state of California have been changed based on a local ordinance. So the benefits representative has asked you to compile a list of the people who are affected. Write a query to display the employee ID and last names of the employees who work in the state of California.

Hint: Use scalar subqueries.

```
SELECT employee_id, last_name
FROM employees e
WHERE ((SELECT location_id
        FROM departments d
        WHERE e.department_id = d.department_id )
IN      (SELECT location_id
        FROM locations l
        WHERE state_province = 'California'));
```

10. Nita, the DBA, wants to remove old information from the database. One of the things she thinks is unnecessary is the old employment records. She has asked you to do the following:

Write a query to delete the oldest JOB_HISTORY row of an employee by looking up the JOB_HISTORY table for the MIN(START_DATE) for the employee. Delete the records of *only* those employees who have changed at least two jobs.

Hint: Use a correlated DELETE command.

```
DELETE FROM job_history JH
WHERE employee_id =
      (SELECT employee_id
       FROM employees E
       WHERE JH.employee_id = E.employee_id
       AND START_DATE = (SELECT MIN(start_date)
                        FROM job_history JH
                        WHERE JH.employee_id =
E.employee_id)
      AND 3 > (SELECT COUNT(*)
               FROM job_history JH
               WHERE JH.employee_id =
E.employee_id)
      GROUP BY EMPLOYEE_ID
      HAVING COUNT(*) >= 2));
```

11. The vice president of Human Resources needs the complete employment records for the annual employee recognition banquet speech. The vice president makes a quick phone call to stop you from following the DBA's orders.

Roll back the transaction.

```
ROLLBACK;
```

12. The sluggish economy is forcing management to take cost reduction actions. The CEO wants to review the highest paid jobs in the company. You are responsible to provide the CEO with a list based on the following specifications:

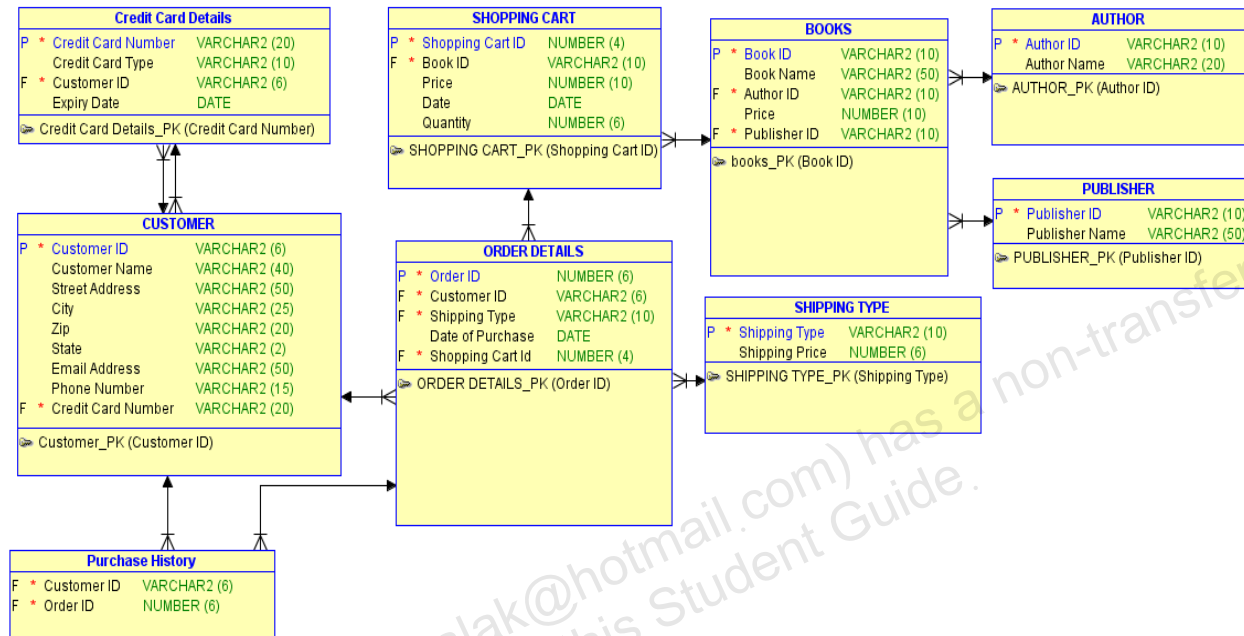
Write a query to display the job IDs of those jobs whose maximum salary is above half the maximum salary in the entire company. Use the WITH clause to write this query. Name the query MAX_SAL_CALC.

```
WITH  
MAX_SAL_CALC AS (SELECT job_title, MAX(salary) AS  
job_total  
FROM employees, jobs  
WHERE employees.job_id = jobs.job_id  
GROUP BY job_title)  
SELECT job_title, job_total  
FROM MAX_SAL_CALC  
WHERE job_total > (SELECT MAX(job_total) * 1/2  
FROM MAX_SAL_CALC)  
ORDER BY job_total DESC;
```

Additional Practices: Case Study

In the case study for the *SQL WORKSHOP I* course, you built a set of database tables for an Online Book Store application. In addition, you inserted, updated, and deleted records in an online book store database and generated a report.

The following is a diagram of the tables and columns that you created for the video application:



Note: First, run the `Online_Book_Store_Drop_Tables.sql` script in the labs folder to drop tables if they already exist. Then run the `Online_Book_Store_Populate.sql` script in the labs folder to create and populate the tables.

1. Verify that the tables were created properly by running a report to show the list of tables and their column definitions.

	TABLE_NAME	COLUMN_NAME	DATA_TYPE	NULLABLE
1	AUTHOR	AUTHOR_ID	VARCHAR2	N
2	AUTHOR	AUTHOR_NAME	VARCHAR2	Y
3	BOOKS	BOOK_ID	VARCHAR2	N
4	BOOKS	BOOK_NAME	VARCHAR2	Y
5	BOOKS	AUTHOR_ID	VARCHAR2	N
6	BOOKS	PRICE	NUMBER	Y
7	BOOKS	PUBLISHER_ID	VARCHAR2	N
8	CREDIT_CARD_DETAILS	CREDIT_CARD_NUMBER	VARCHAR2	N
9	CREDIT_CARD_DETAILS	CREDIT_CARD_TYPE	VARCHAR2	Y
10	CREDIT_CARD_DETAILS	EXPIRY_DATE	DATE	Y
11	CUSTOMER	CUSTOMER_ID	VARCHAR2	N
12	CUSTOMER	CUSTOMER_NAME	VARCHAR2	Y
13	CUSTOMER	STREET_ADDRESS	VARCHAR2	Y
14	CUSTOMER	CITY	VARCHAR2	Y
15	CUSTOMER	PHONE_NUMBER	VARCHAR2	Y
16	CUSTOMER	CREDIT_CARD_NUMBER	VARCHAR2	N
17	ORDER_DETAILS	ORDER_ID	VARCHAR2	N
18	ORDER_DETAILS	CUSTOMER_ID	VARCHAR2	Y
19	ORDER_DETAILS	SHIPPING_TYPE	VARCHAR2	N
20	ORDER_DETAILS	DATE_OF_PURCHASE	DATE	Y
21	ORDER_DETAILS	SHOPPING_CART_ID	VARCHAR2	N
22	PUBLISHER	PUBLISHER_ID	VARCHAR2	N
23	PUBLISHER	PUBLISHER_NAME	VARCHAR2	Y
24	PURCHASE_HISTORY	CUSTOMER_ID	VARCHAR2	Y
25	PURCHASE_HISTORY	ORDER_ID	VARCHAR2	N
26	SHIPPING_TYPE	SHIPPING_TYPE	VARCHAR2	N
27	SHIPPING_TYPE	SHIPPING_PRICE	NUMBER	Y
28	SHOPPING_CART	SHOPPING_CART_ID	VARCHAR2	N
29	SHOPPING_CART	BOOK_ID	VARCHAR2	N
30	SHOPPING_CART	PRICE	NUMBER	Y
31	SHOPPING_CART	SHOPPING_CART_DATE	DATE	Y
32	SHOPPING_CART	QUANTITY	NUMBER	Y

- Verify the existence of the ORDER_ID_SEQ sequence in the data dictionary.

	SEQUENCE_NAME
1	DEPARTMENTS_SEQ
2	EMPLOYEES_SEQ
3	LOCATIONS_SEQ
4	ORDER_ID_SEQ

- You want to create some users who have access only to their purchase history. Create a user called Carmen and grant her the privilege to select from the PURCHASE_HISTORY table.
Note: Make sure to prefix the username with your database account. For example, if you are the user oraxx, create a user called oraxx_Carmen.
- Add an edition column (varchar2 (6)) to the BOOKS table to store the book edition information.
- Add a CREDIT_CARD_TYPE table to store CREDIT_CARD_TYPE and CREDIT_CARD_DESCRIPTION. The table has a foreign key with the CREDIT_CARD_TYPE column in the CREDIT_CARD_DETAILS table.
- Select all the tables from the data dictionary.
- Create a SHOPPING_HISTORY table to store the details of purchase history of the customers.
(Hint: You can copy the PURCHASE_HISTORY table.)
- Display the customer details of the first ten customers who have placed orders in the last month. Order the records based on the customer ID.

	CUSTOMER_ID	ORDER_ID	DATE_OF_PURCHASE	CUSTOMER_NAME
1	CN0001	0D0001	12-JUN-01	VelasquezCarmen
2	CN0003	0D0003	31-JUL-05	Nagayama Midori
3	CN0004	0D0004	14-AUG-06	Quick-To-See Mark
4	CN0009	0D0009	25-NOV-09	Catchpole Antoinette

- Show a list of customers who have placed an order more than one time.

	CUSTOMER_ID	CUSTOMER_NAME
1	CN0001	VelasquezCarmen
2	CN0003	Nagayama Midori
3	CN0004	Quick-To-See Mark
4	CN0009	Catchpole Antoinette

Additional Practices Solution: Case Study

Solution

First, run the `Online_Book_Store_Drop_Tables.sql` script in the labs folder to drop tables if they already exist. Then run the `Online_Book_Store_Populate.sql` script in the labs folder to create and populate the tables.

1. Verify that the tables were created properly by running a report to show the list of tables and their column definitions.

```
SELECT table_name,column_name,data_type,nullable
FROM user_tab_columns
WHERE table_name
IN('CUSTOMER','CREDIT_CARD_DETAILS','SHOPPING_CART',
'ORDER_DETAILS','BOOKS','AUTHOR','PUBLISHER','SHIPPING_TYPE',
'PURCHASE_HISTORY');
```

2. Verify the existence of the `ORDER_ID_SEQ` sequences in the data dictionary.

```
SELECT sequence_name FROM user_sequences;
```

3. You want to create some users who have access only to their purchase history. Create a user called Carmen and grant her the privilege to select from the `PURCHASE_HISTORY` table.

Note: Make sure to prefix the username with your database account. For example, if you are the user `oraxx`, create a user called `oraxx_Carmen`.

```
CREATE USER oraxx_carmen IDENTIFIED BY oracle ;
GRANT select ON purchase_history TO oraxx_carmen;
```

4. Add an edition column (`varchar2 (6)`) to the `BOOKS` table to store the book edition information.

```
ALTER TABLE books ADD(edition VARCHAR2(6));
```

5. Add a `CREDIT_CARD_TYPE` table to store `CREDIT_CARD_TYPE` and `CREDIT_CARD_DESCRIPTION`. The table has a foreign key with the `CREDIT_CARD_TYPE` column in the `CREDIT_CARD_DETAILS` table.

```
CREATE TABLE CREDIT_CARD_TYPE
( CREDIT_CARD_TYPE VARCHAR2(10) NOT NULL ENABLE,
  CREDIT_CARD_DESCRIPTION VARCHAR2(4000 BYTE),
  CONSTRAINT CREDIT_CARD_TYPE_PK PRIMARY KEY
  (CREDIT_CARD_TYPE) )
;
```


6. Select all the tables from the data dictionary.

```
SELECT table_name FROM user_tables order by table_name;
```

7. Create a SHOPPING_HISTORY table to store the details of a purchase history of customers. (Hint: You can copy the PURCHASE_HISTORY table.)

```
CREATE TABLE shopping_history as select * from purchase_history  
where '1' = '1';
```

8. Display the customer details of the first ten customers who have placed orders in the last month. Order the records based on the customer ID.

```
SELECT o.CUSTOMER_ID, o.ORDER_ID, o.DATE_OF_PURCHASE,  
c.CUSTOMER_NAME  
FROM ORDER_DETAILS o JOIN PURCHASE_HISTORY p  
ON o.CUSTOMER_ID = p.CUSTOMER_ID JOIN CUSTOMER c  
ON o.CUSTOMER_ID= c.CUSTOMER_ID  
AND rownum < 10  
ORDER BY CUSTOMER_ID;
```

9. Show a list of customers who have placed an order more than one time.

```
SELECT customer_id, customer_name FROM customer c  
WHERE 1 <= (select count(*) from purchase_history where  
customer_id = c.customer_id);
```

YASSINE MAZOUZ (z.m.malak@hotmail.com) has a non-transferable
license to use this Student Guide.